

## Real-Time Processing with Kafka

---

# Learning Goals



## Learning Goals

---



- Summarize the motivation behind Kafka Streams
- Define core components of Kafka Streams
- Summarize the life of a message in Kafka Streams
- Build a simple producer and consumer application

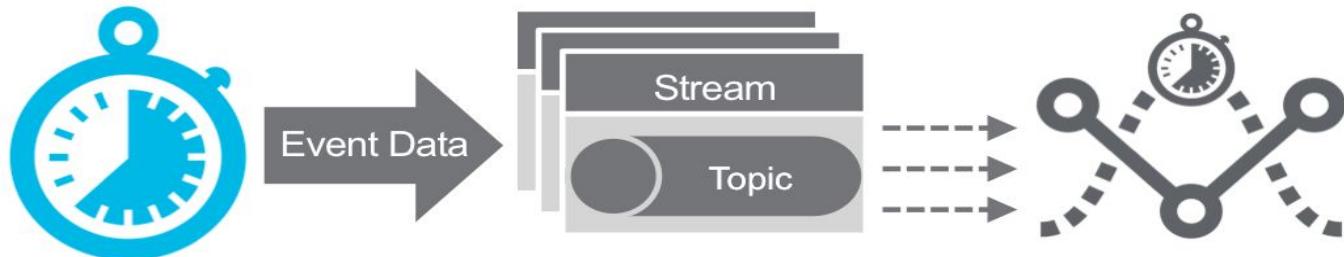
## Prerequisites

---

- Basic understanding of big data concepts
- Basic understanding of the platform
- Basic understanding of application development principles

---

## Many Big Data sources are Event Oriented



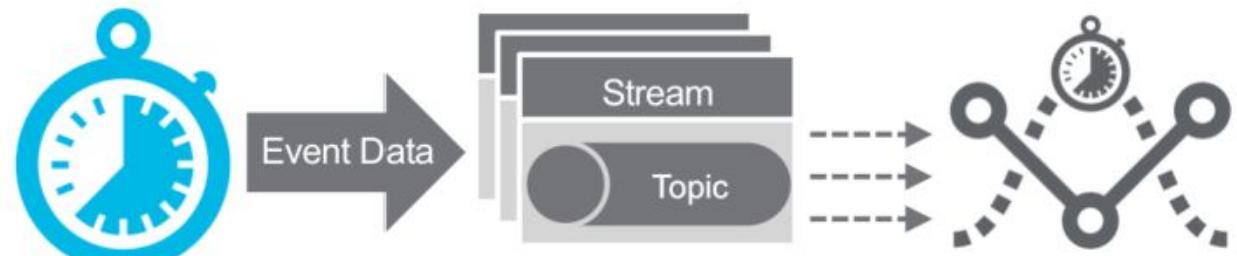
**Trigger Events:**

- Stock Prices
- User Activity
- Sensor Data

Real-Time Analytics

# What is Kafka?

Many Big Data sources are Event Oriented



## Trigger Events:

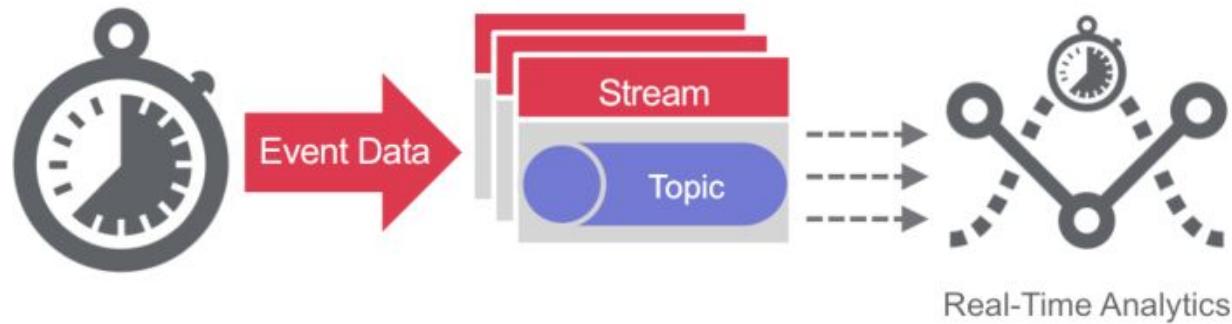
- Stock Prices
- User Activity
- Sensor Data

## Real-Time Analytics

# Why Kafka?

## Today's Applications need to:

- process high-velocity data as soon as possible
- handle high-volume workloads

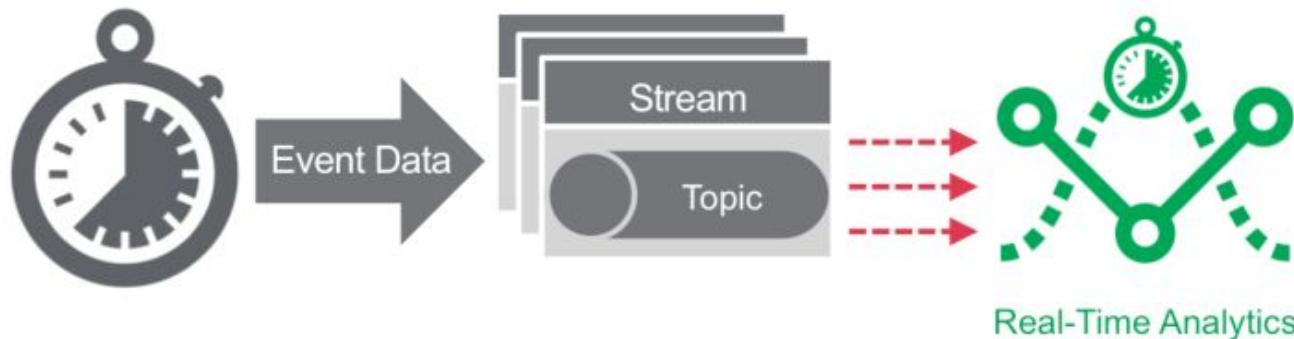


# Why Kafka?

---

**Today's applications need to:**

- process high-velocity data as soon as possible
- handle high-volume workloads
- provide real-time results with extremely low latency



## Diverse Data Sources

---



Social Media Feeds



Financial Transactions



Geotagging Data



Internet of Things



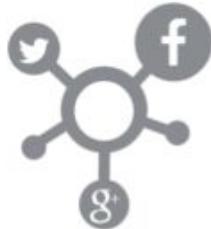
Web Browser Clickstreams



Application Metrics

## Diverse Data Sources

---



Social Media Feeds



Financial Transactions



Geotagging Data



Internet of Things



Web Browser Clickstreams



Application Metrics

## Diverse Data Sources

---



Social Media Feeds



Financial Transactions



Geotagging Data



Internet of Things



Web Browser Clickstreams



Application Metrics

## Diverse Data Sources

---



Social Media Feeds



Financial Transactions



Geotagging Data



Internet of Things

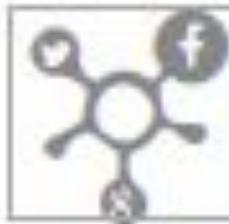


Web Browser Clickstreams



Application Metrics

## Diverse Data Sources



Social Media Feeds



Financial Transactions



Geolocation Data



Internet of Things



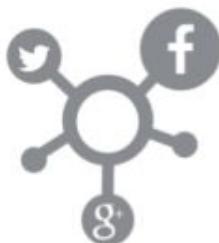
Web Browser Clickstreams



Application Metrics

## Diverse Data Sources

---



Social Media Feeds



Financial Transactions



Geotagging Data



Internet of Things



Web Browser Clickstreams



Application Metrics

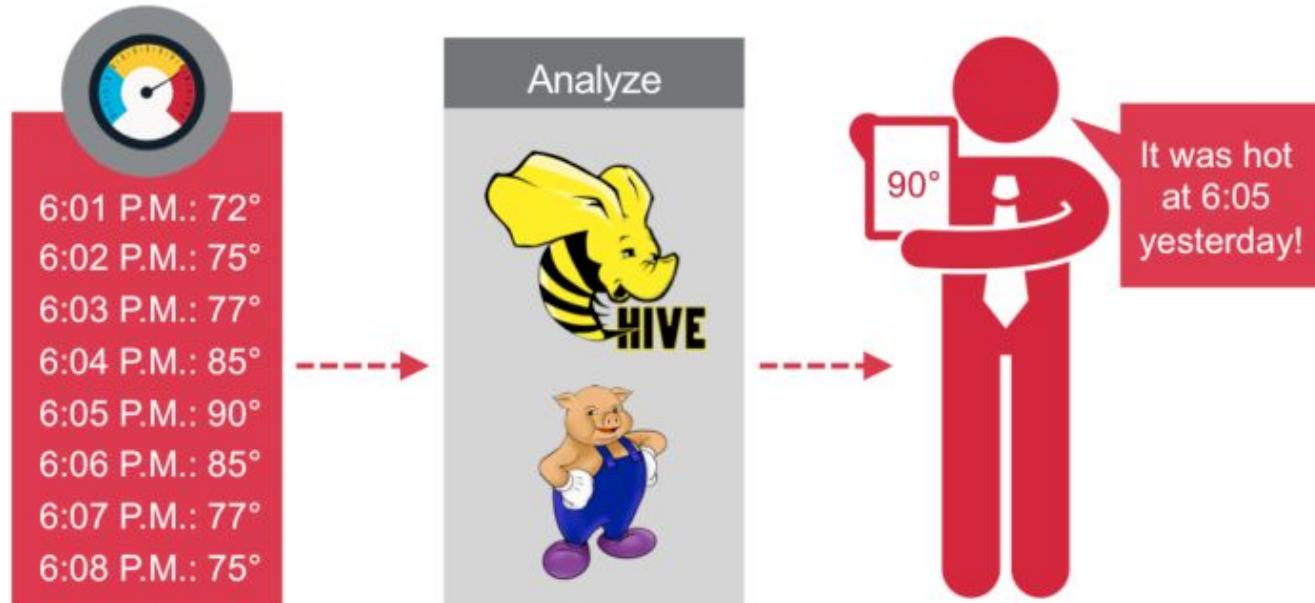
## Analyze Data

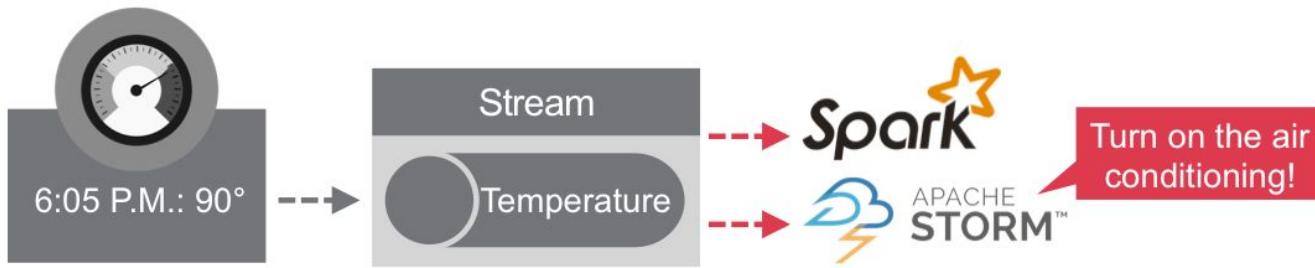
---

What if you need to analyze data as it arrives?



# Batch Processing with HDFS





## Organize Data

---

What if you need to organize data as it arrives?

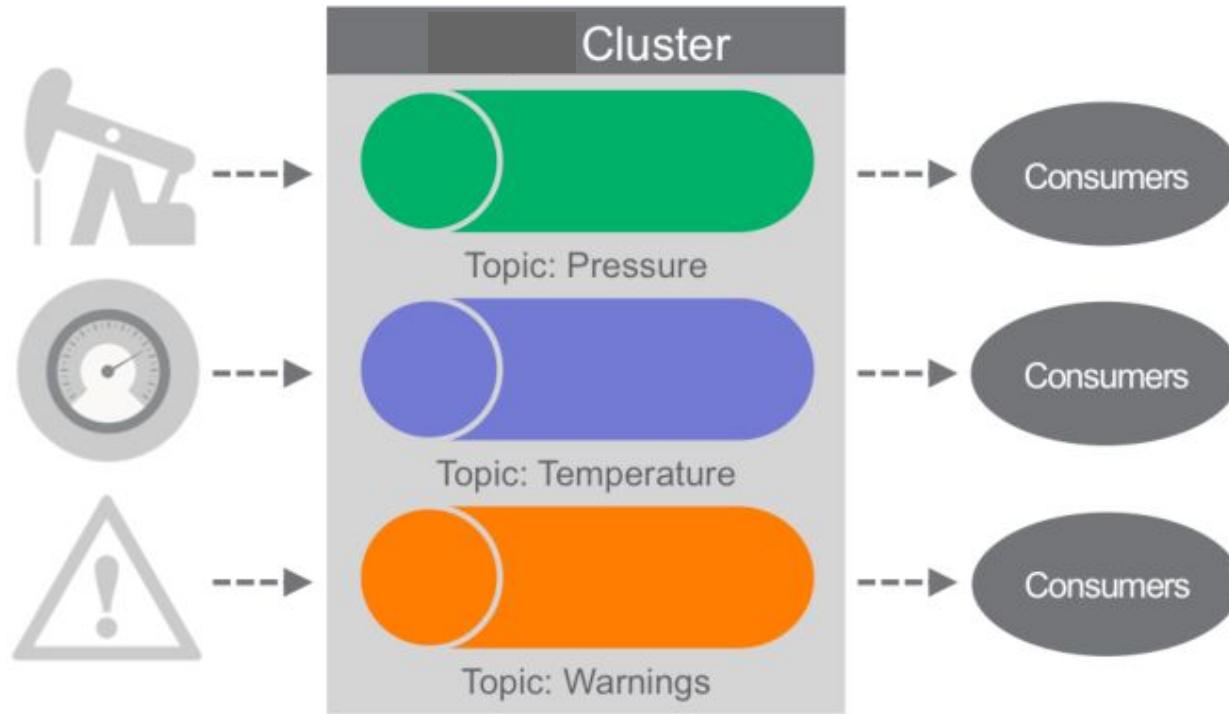


# Many Sources of Data



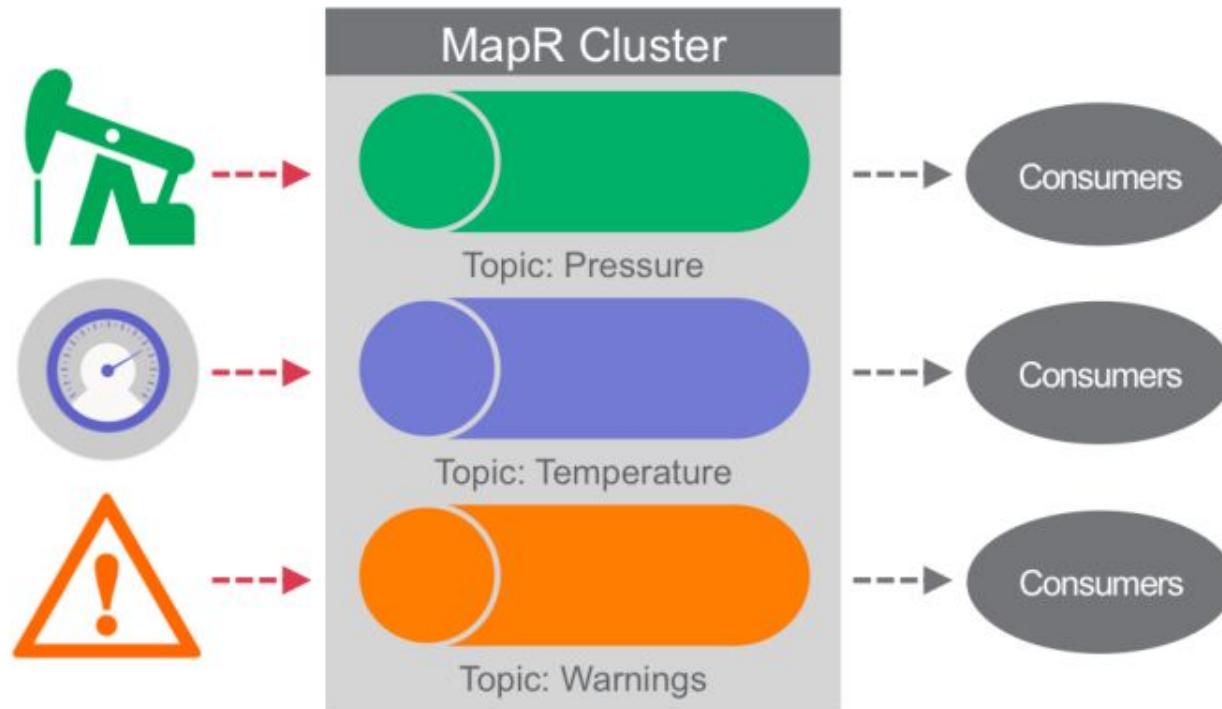
# Organize Data into Topics with Kafka Streams

Topics organize events into categories



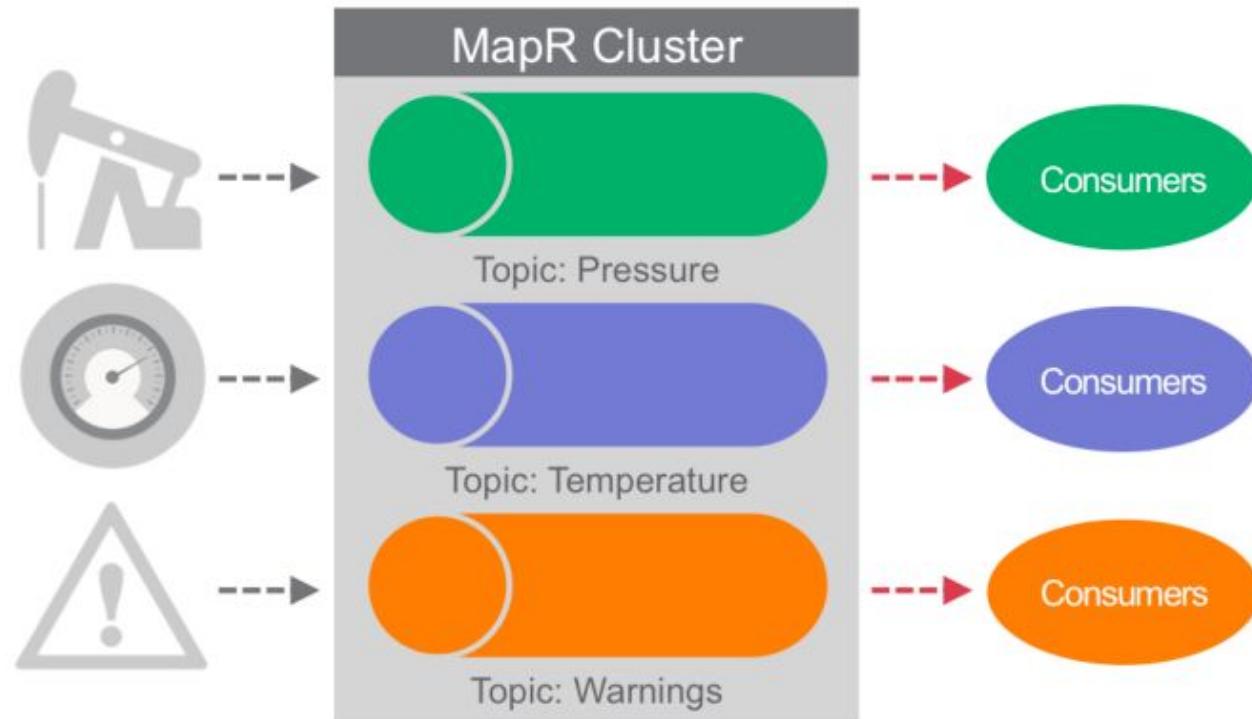
## Organize Data into Topics with MapR Streams

Producers publish to topics



# Organize Data into Topics with MapR Streams

Consumers subscribe to topics



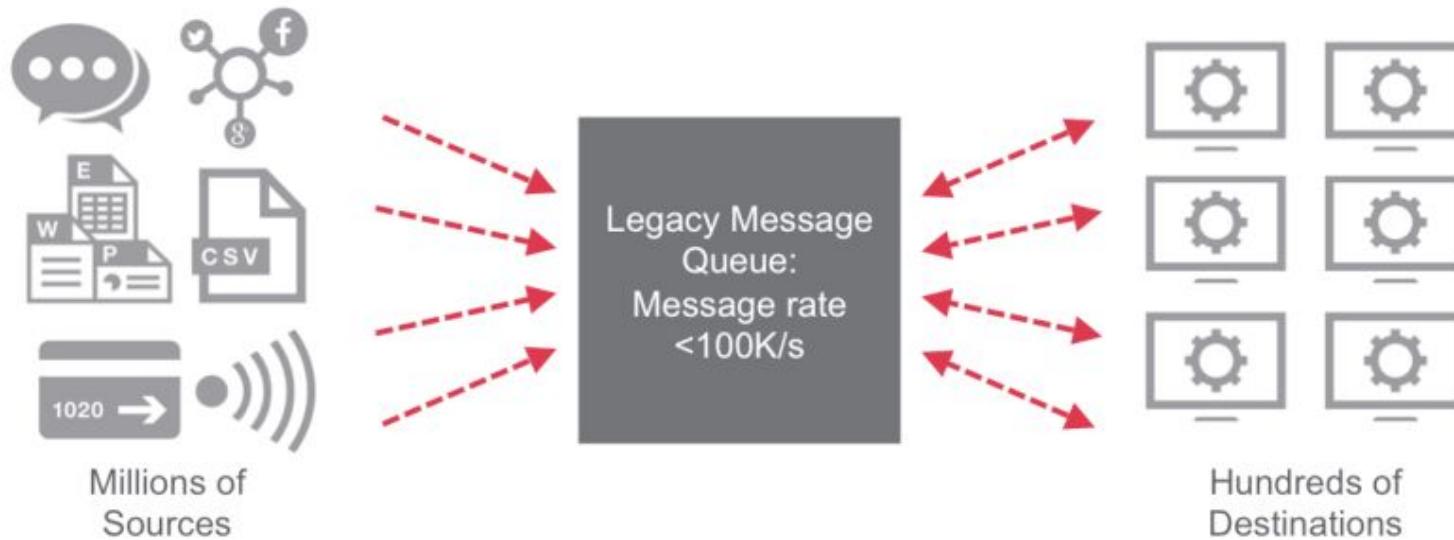
## Process High Volume of Data

---

What if you need to process a high volume of data as it arrives?

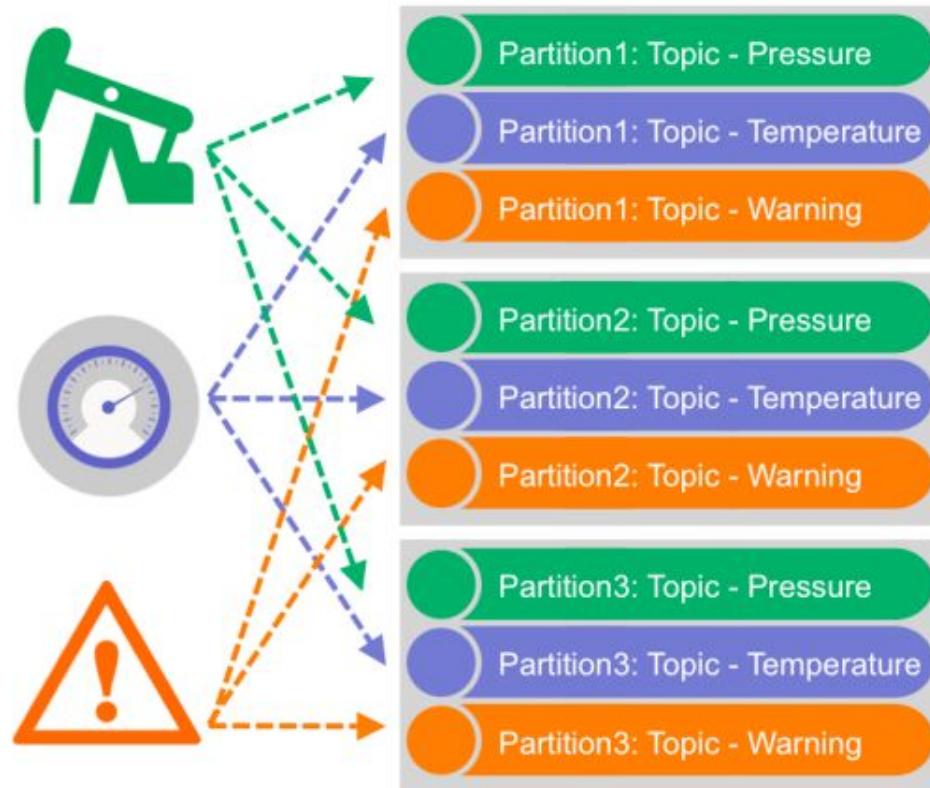


# Legacy Messaging

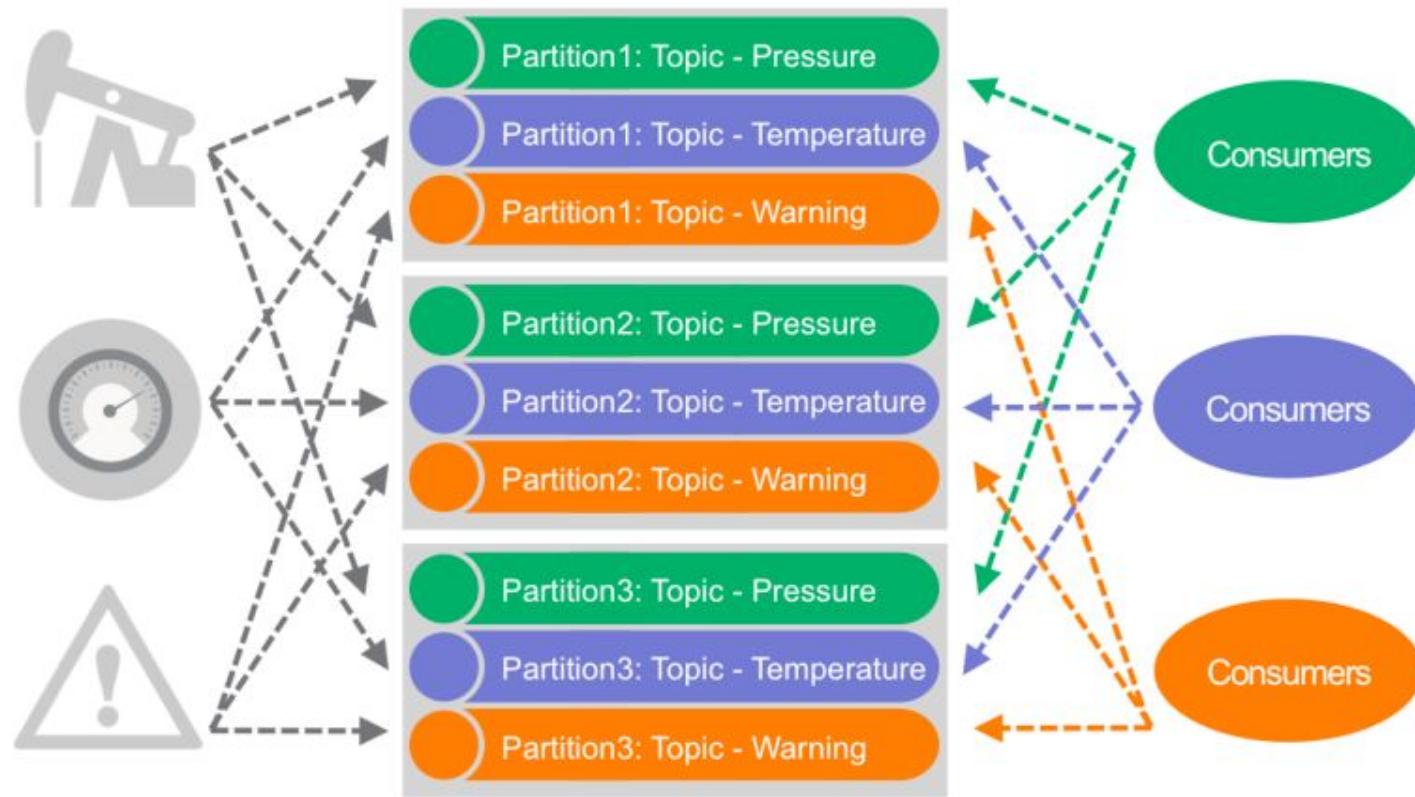




# Scalable Messaging with MapR Streams



# Scalable Messaging with MapR Streams



## Message Recovery

---

What if you need to recover messages in case of server failure?



## Lack of Replication



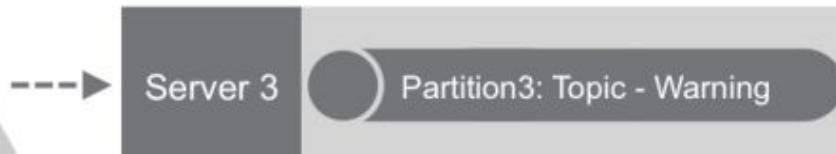
Producer



Producer



Producer



## Partitions are Replicated for Fault Tolerance



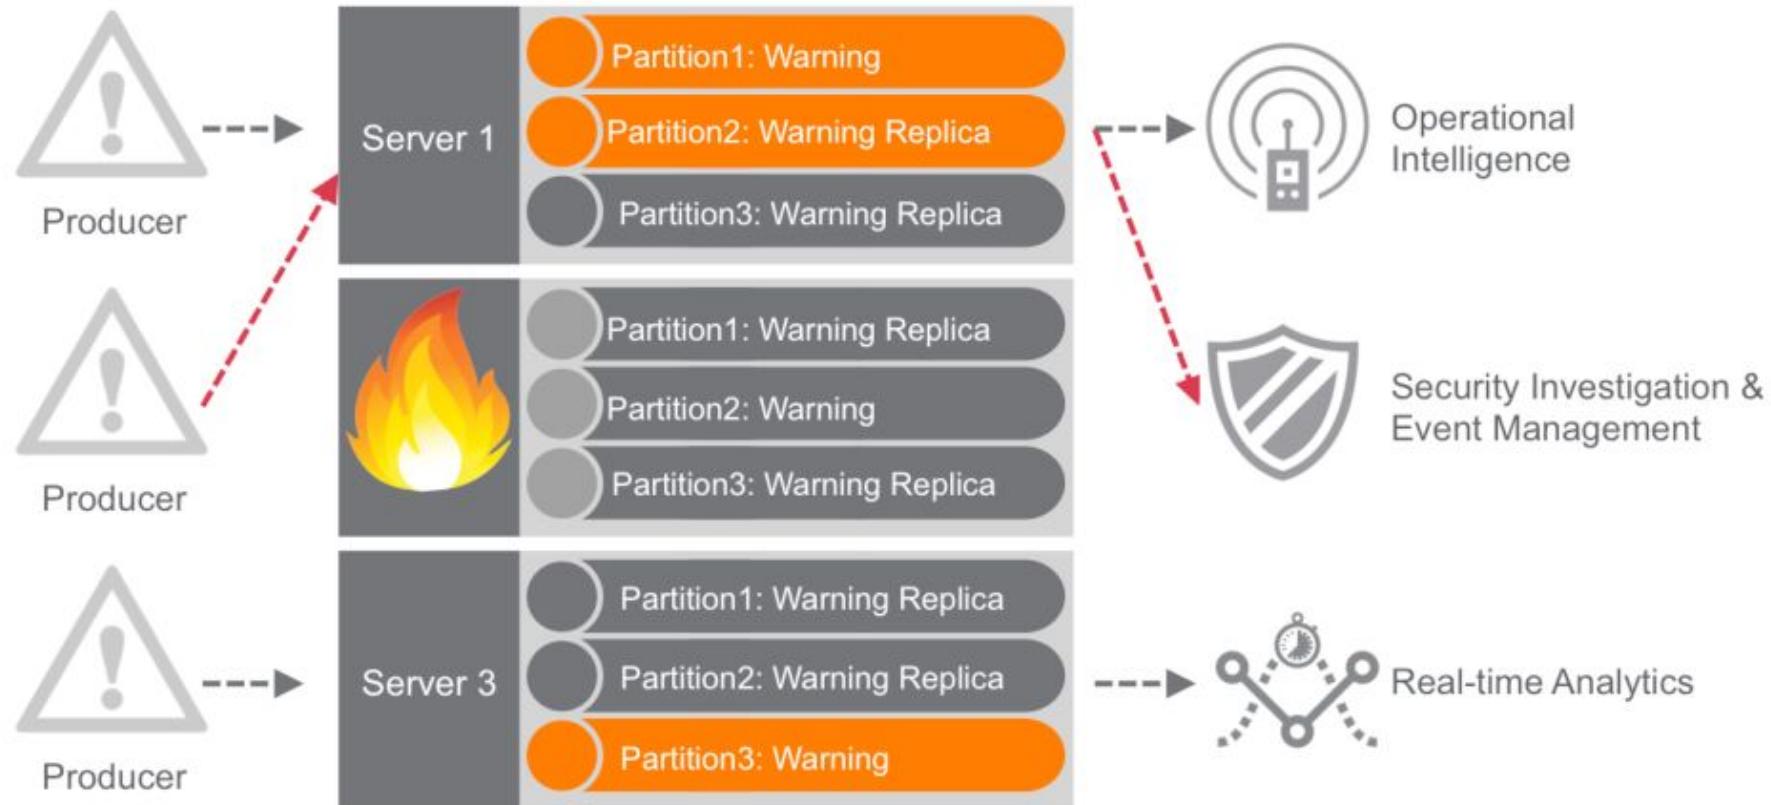
# Partitions are Replicated for Fault Tolerance



# Partitions are Replicated for Fault Tolerance



# Partitions are Replicated for Fault tolerance



## Real-time Access

---

What if you need real-time access to live data distributed across multiple clusters and multiple data centers?



## Lack of Global Replication

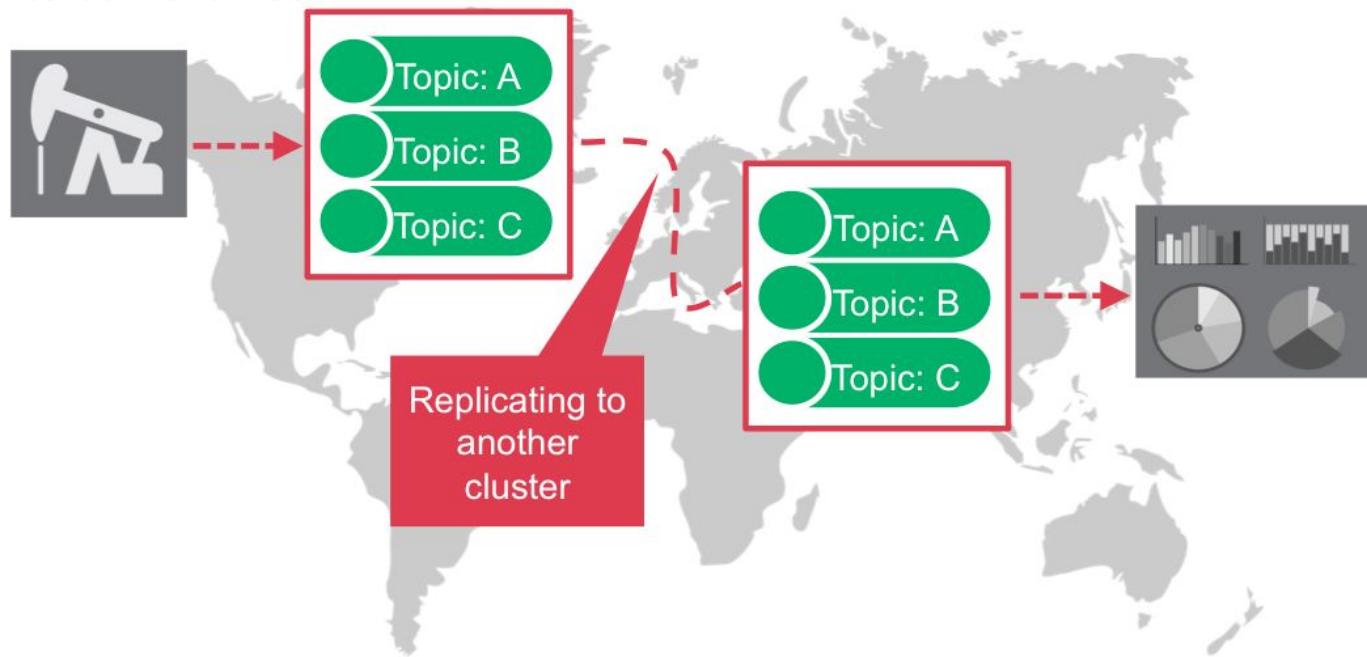
---



# Streams and Replication

## Streams:

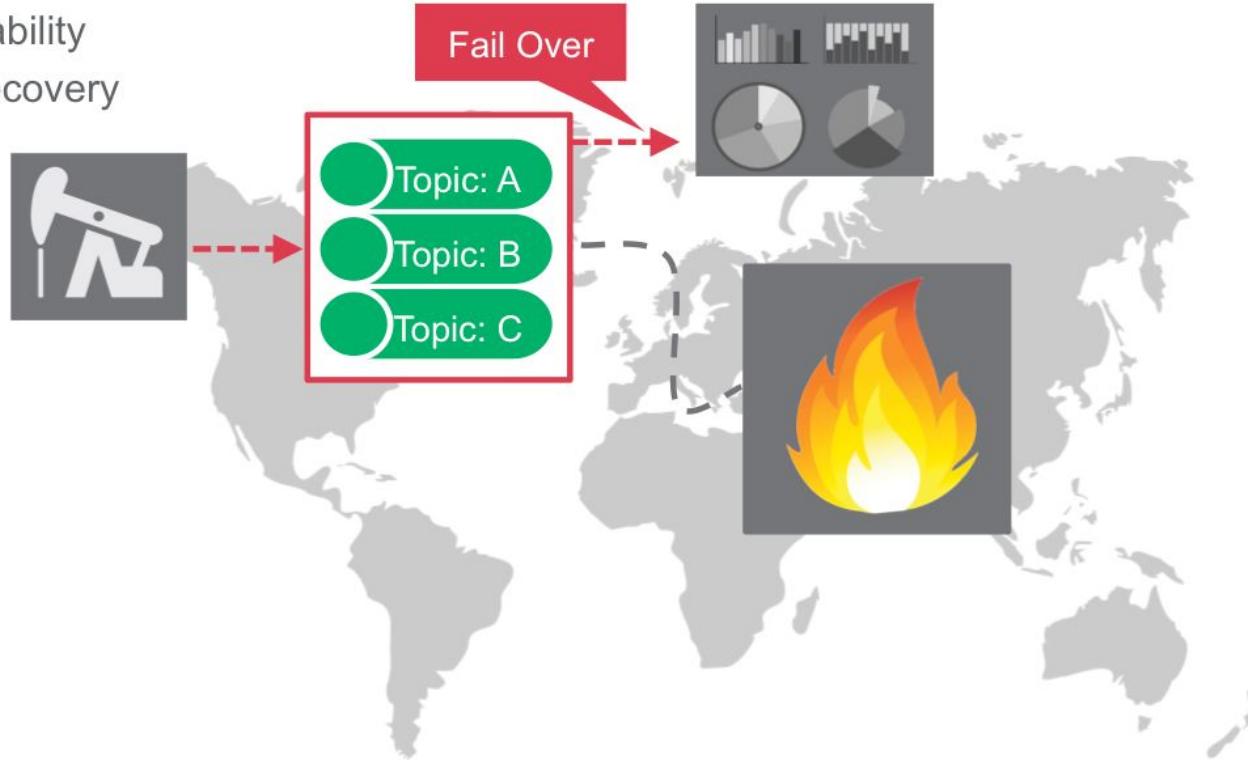
- are a collection of topics
- can be replicated worldwide



# Streams and Replication

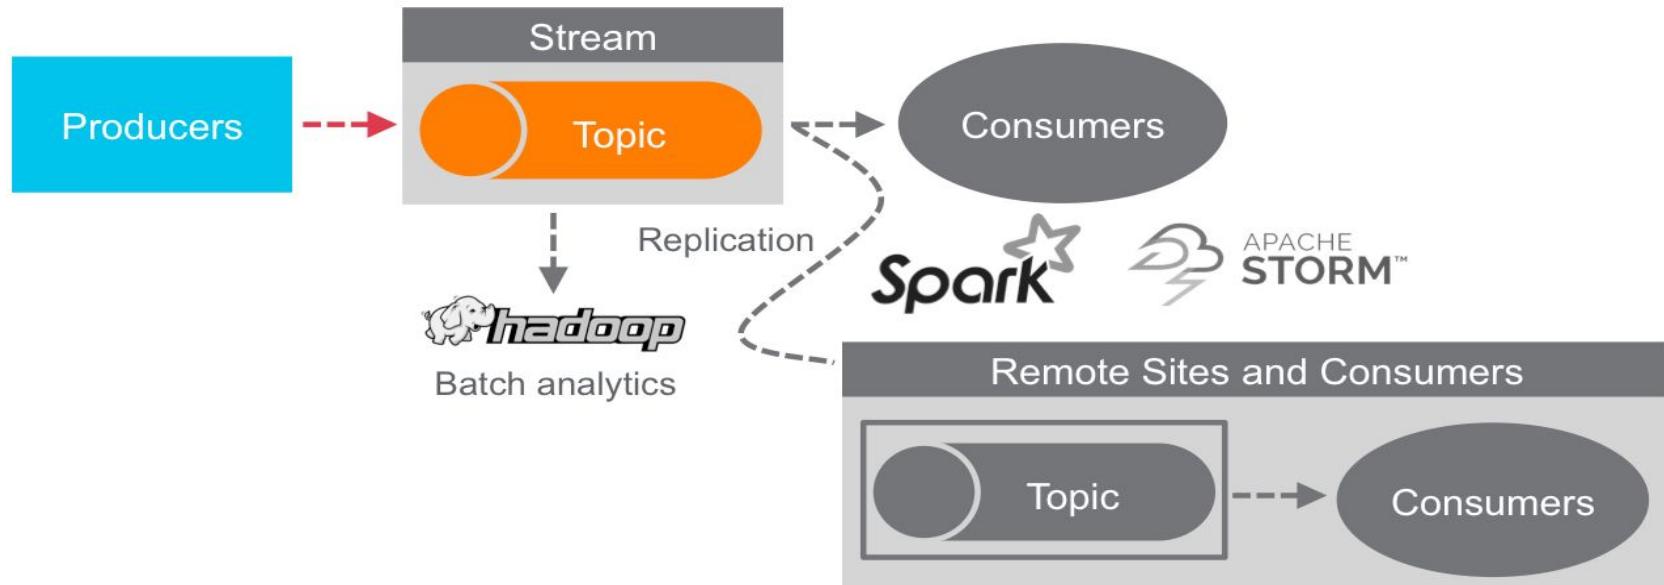
## Streams:

- high availability
- disaster recovery

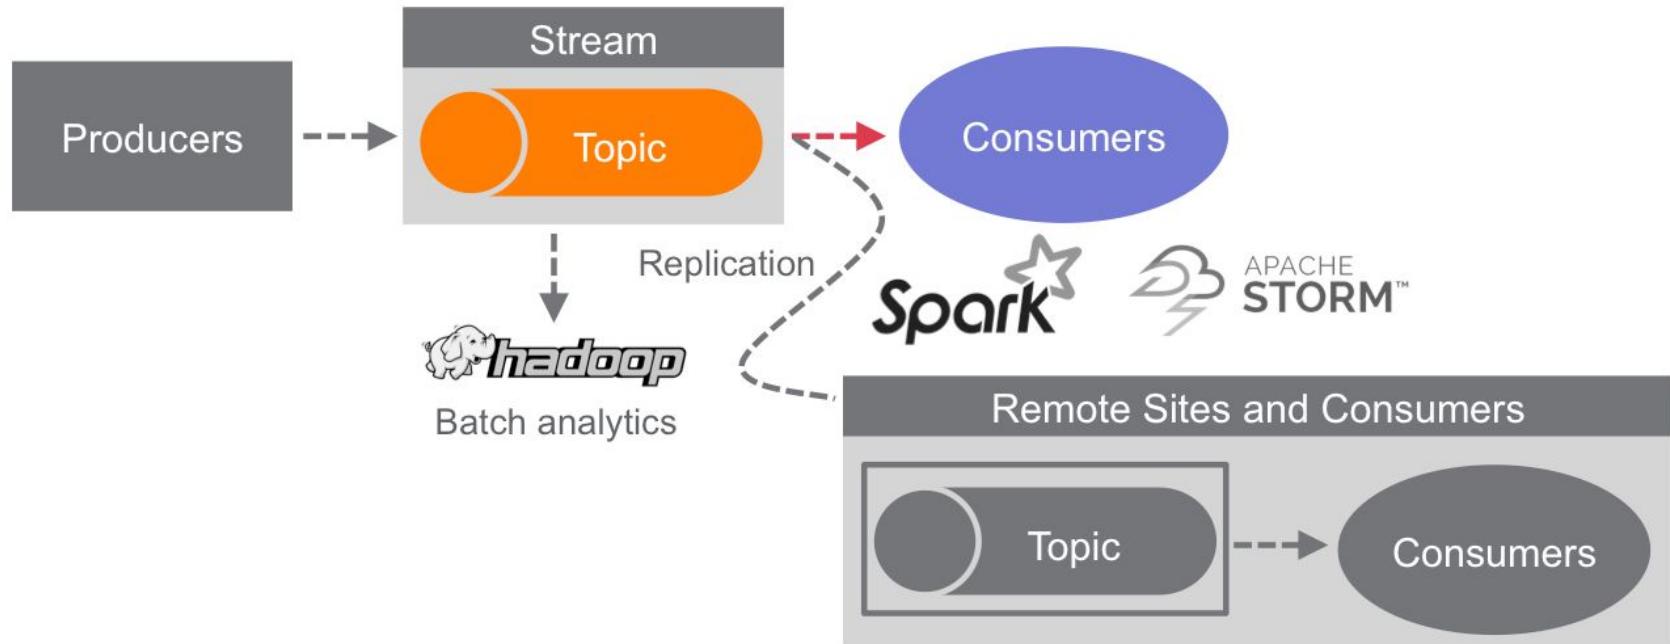


# Kafka: A Global Pub-Sub System for Big Data

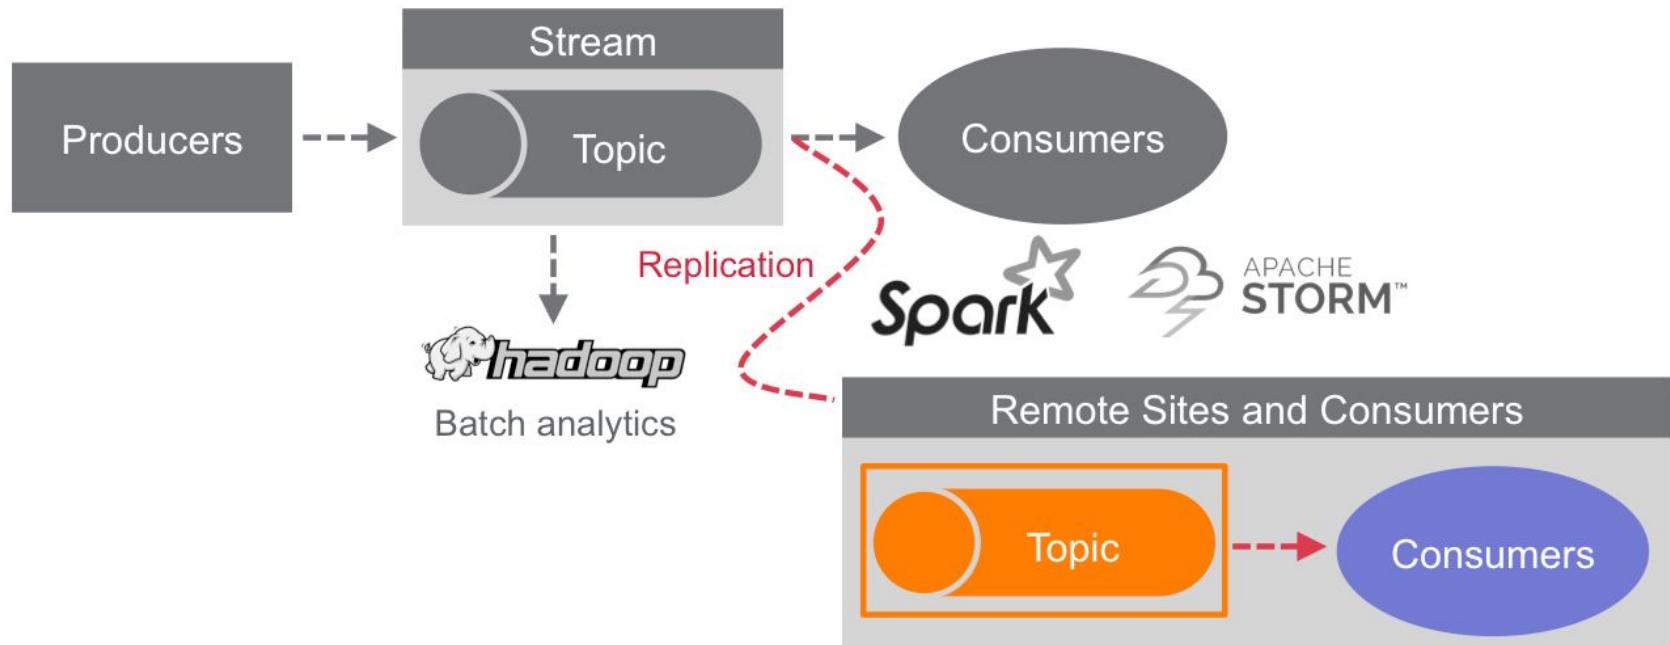
Producers publish billions of events/sec



## Reliable delivery to all consumers

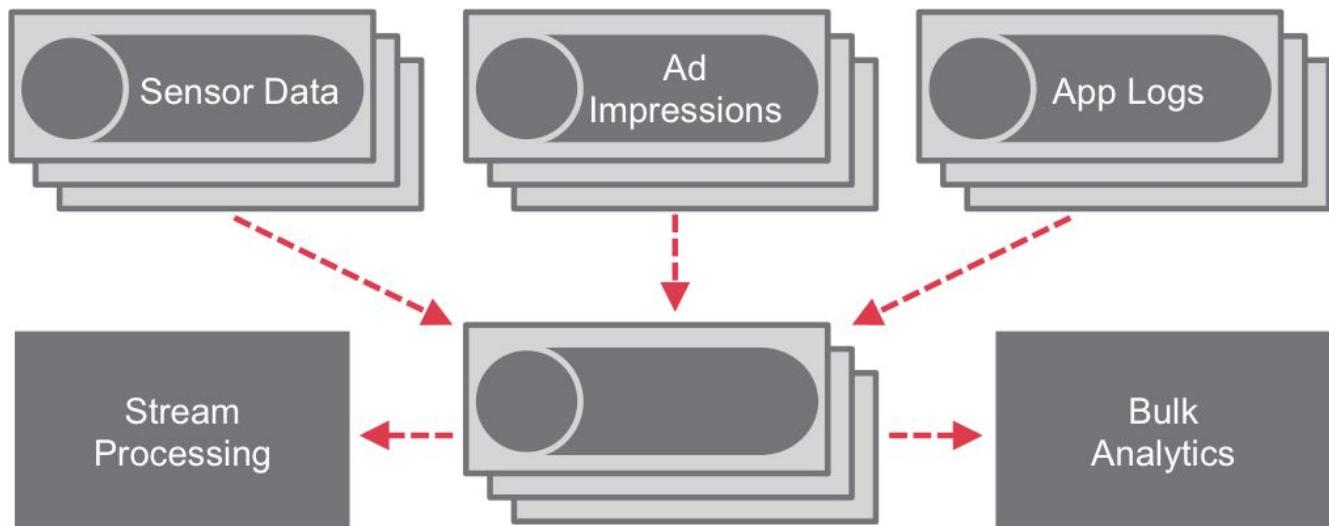


Tie together geo-dispersed clusters worldwide



- 
- Combines:
    - Massively scalable messaging
    - Consistent, reliable delivery
    - Real-time analytics
    - Security and fault-tolerance
  - Standard real-time API (Kafka).
  - OJAI API – direct data access from analytics frameworks

- Scalable to handle Internet of Things
- Global producers and consumers
- Secure and fault-tolerant
- Bulk and real-time analytics





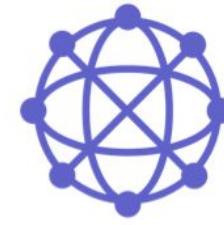
Health Care



Advertising



Credit Card Security



Internet of Things

# Use Case: Streaming System of Record for Healthcare

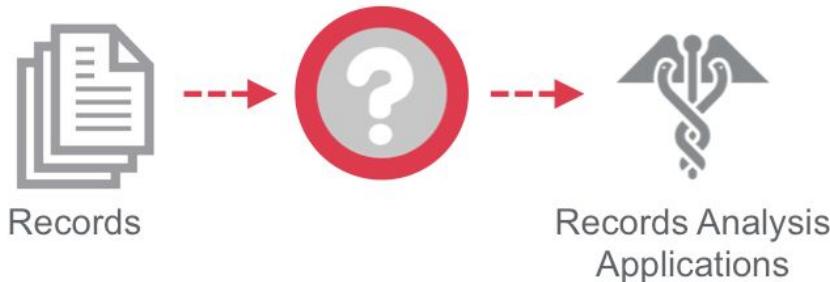
---

## Objective:

- Build a flexible, secure healthcare database

## Challenges:

- Many different data models
- Security and privacy issues
- HIPAA compliance



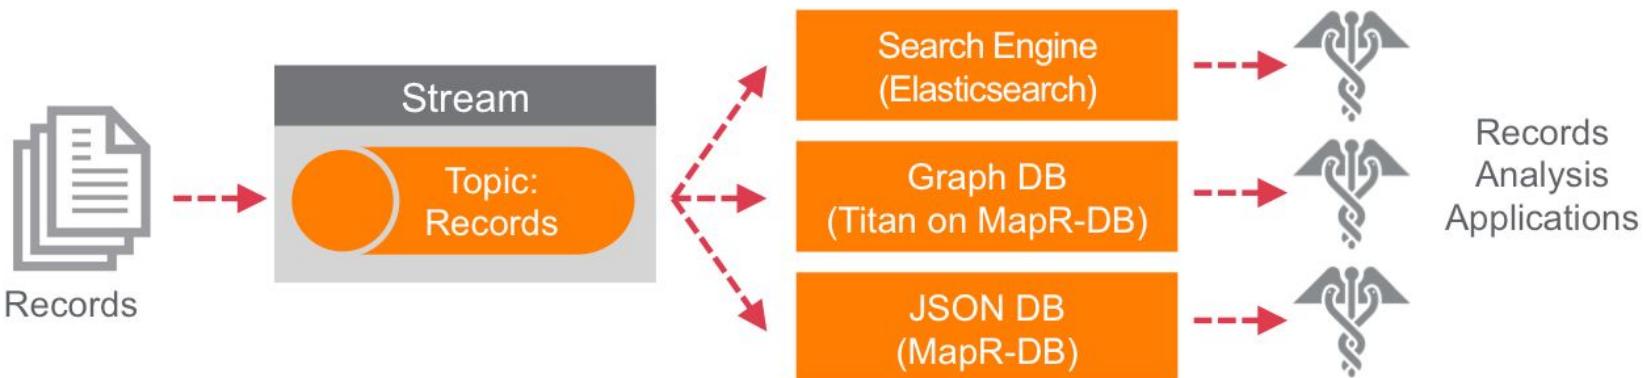
# Use Case: Streaming System of Record for Healthcare

## Solution:

- Streaming system of record
  - secure
  - immutable
  - rewritable

## Auditable

- Materialized views continuously computed
- Selective cross data center replication



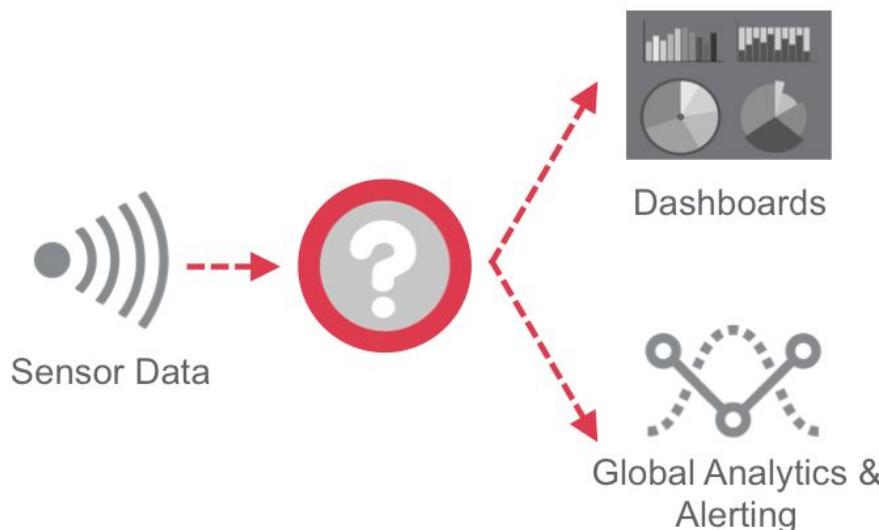
# Use Case: Monitoring the Internet of Things in Real-Time

## Objective:

- Monitor oil rig sensor data and create real-time alerts

## Challenges:

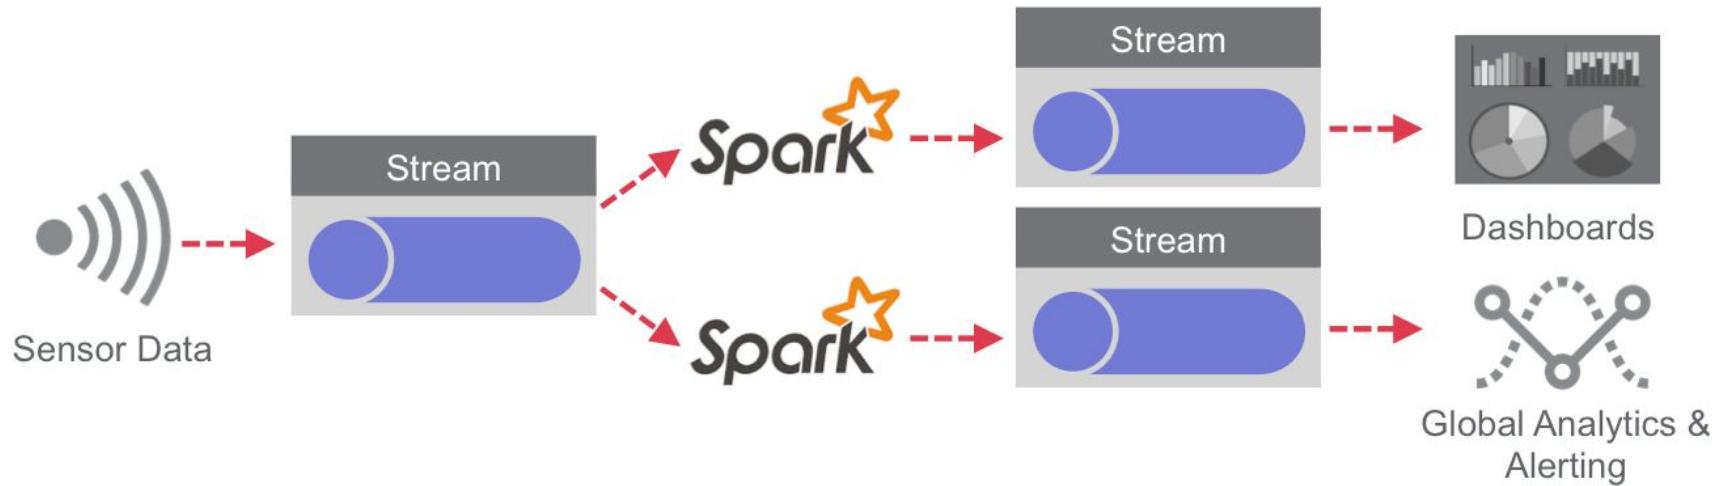
- Disperse, global data sources
- Need alerts in real-time
- Need to audit data



# Use Case: Monitoring the Internet of Things in Real-Time

## Solution:

- Synchronize global servers
- Streaming analytics
- Replicated streams



# Use Case: Global Analytics in Targeted Advertising

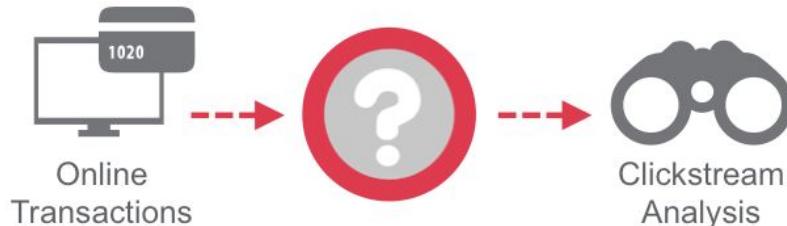
---

## Objective:

- Optimize global ad market with real-time analytics

## Challenges:

- Data consistency
- Global insights
- Complex pipeline

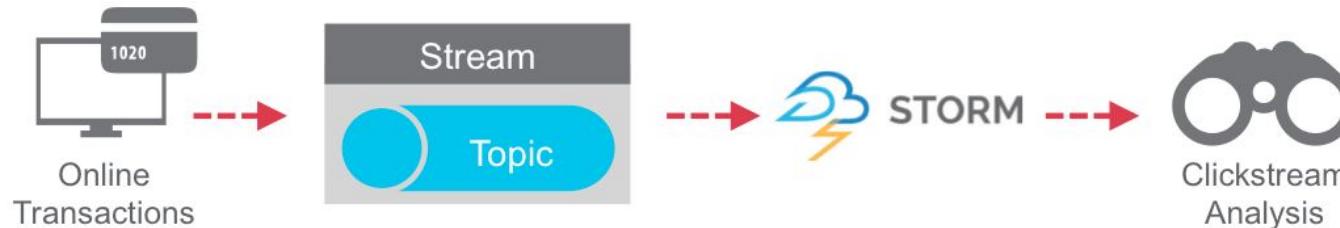


# Use Case: Global Analytics in Targeted Advertising

---

## Solution:

- Real-time analytics
- Global message streams
- Streaming ETL



# Use Case: Intelligent Credit Card Processing

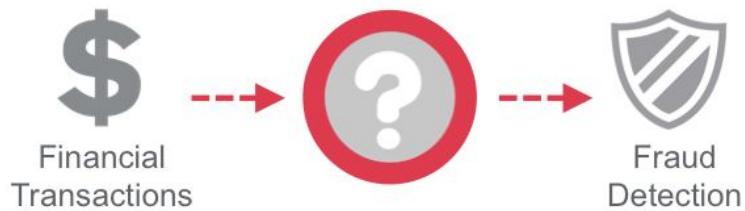
---

## Objective:

- Build reliable, real-time service bus for financial transactions

## Challenges:

- Accurate data
- Sensitive data
- Ad-hoc querying



# Use Case: Intelligent Credit Card Processing

---

## Solution:

- Reliable, replicated streams
- Security settings





# Goals

## Core Components: Message

---

- Messages are key/value pairs
- Keys are optional
- Values can be any type of data

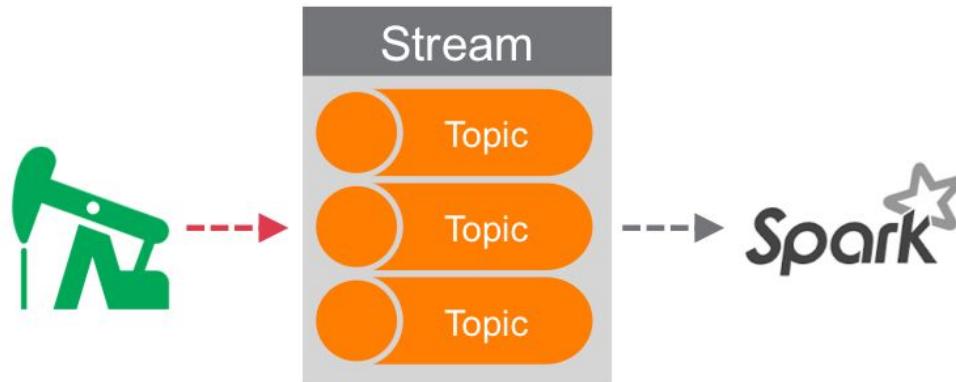
key	value
	{"event_ID":122, "hcid":63, "timestamp": "2015-01-01 00:00:02", "production":47, "sensor_readings": {"pressure":107, "temperature":30, "oil_percentage":18}, "injection_vol": 15367}

# Core Components: Producers and Consumers

---

## Producers

- publish to topics
- example: networked sensors on an oil rig

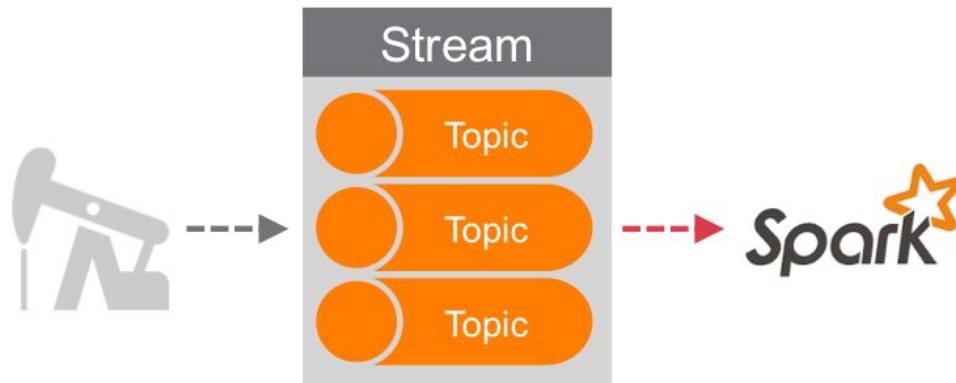


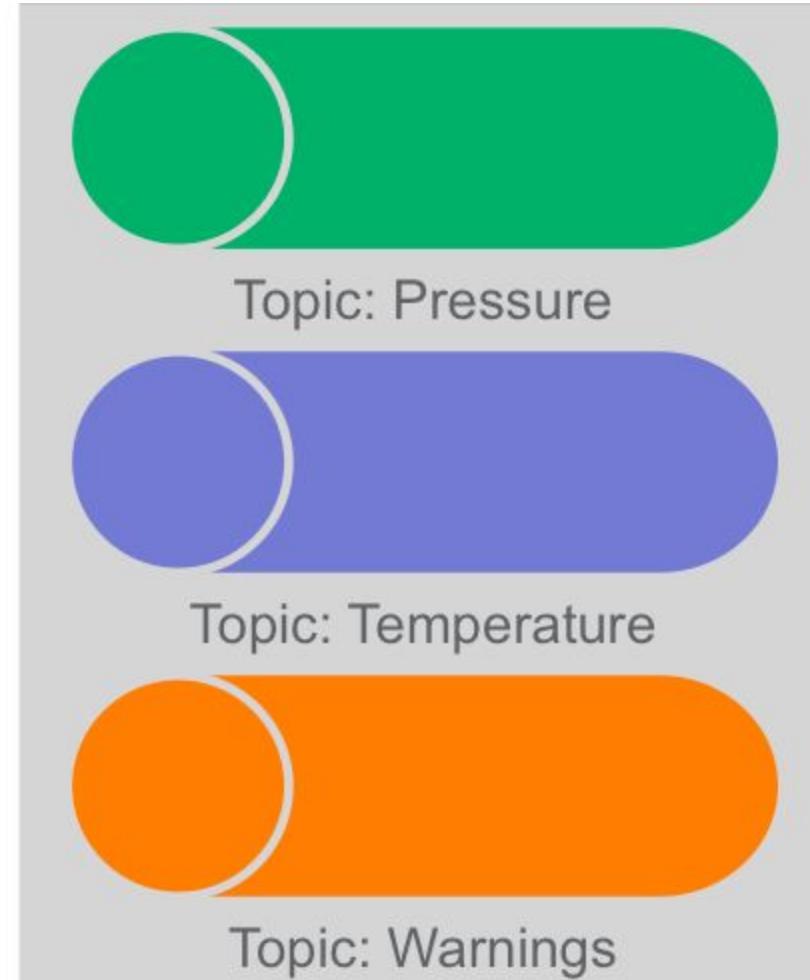
# Core Components: Producers and Consumers

---

## Consumers

- subscribe to topics
- example: applications using Apache Spark







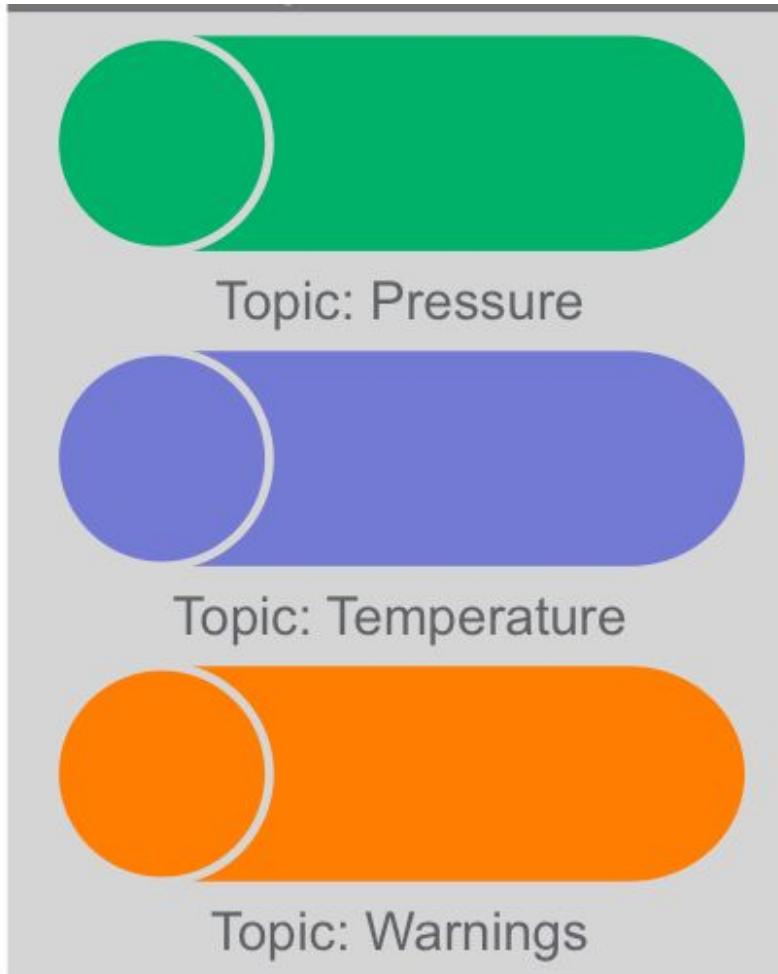
Topic: Pressure



Topic: Temperature



Topic: Warnings



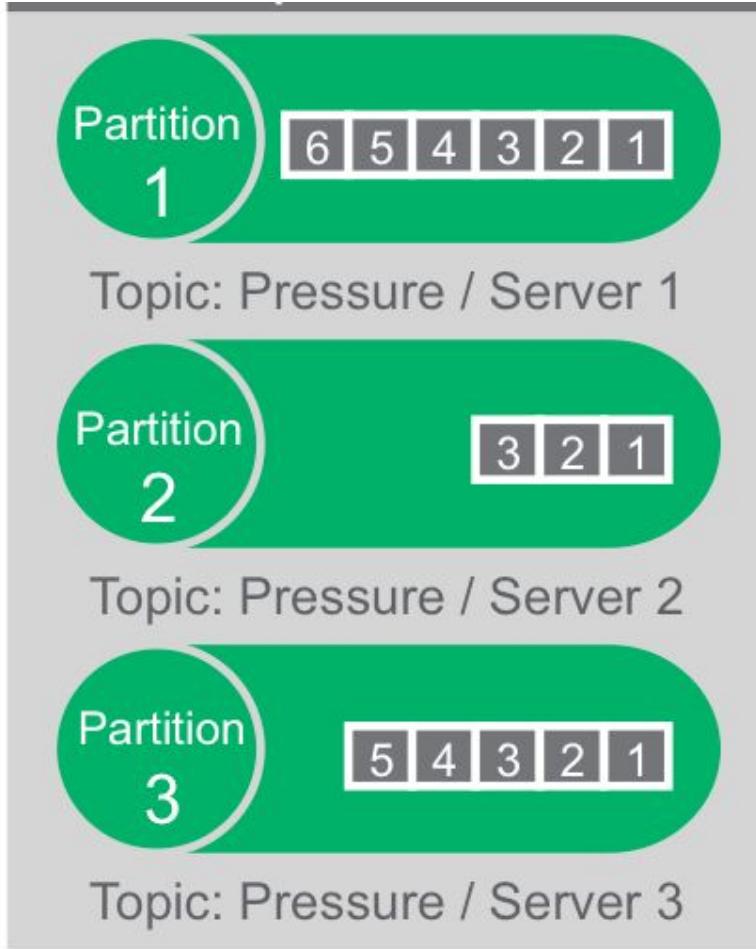
Consumers

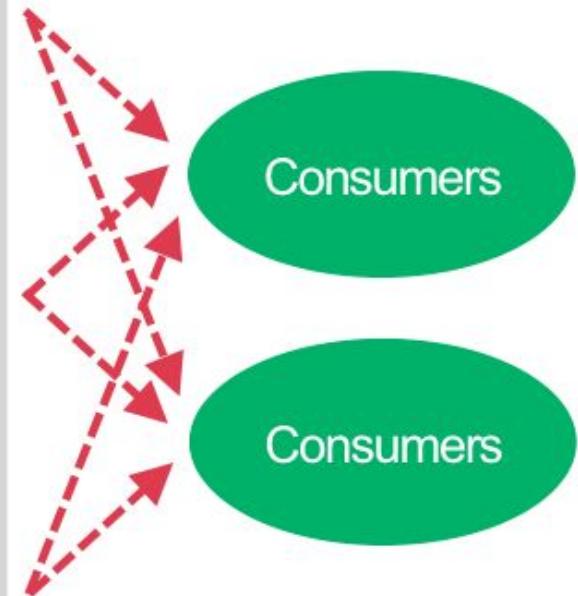


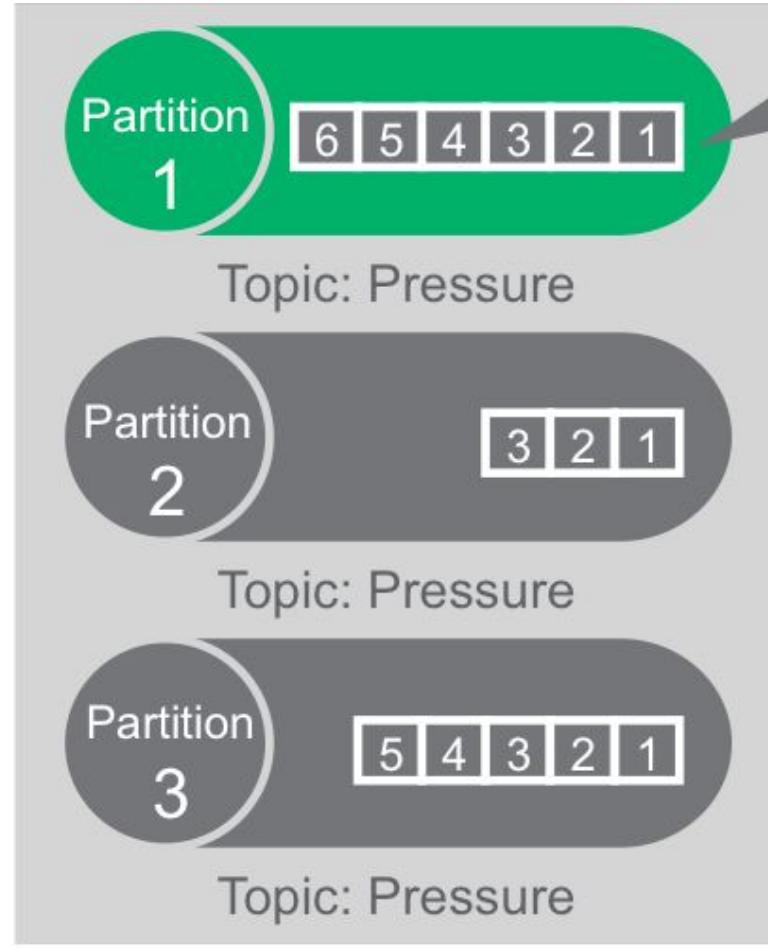
Consumers



Consumers





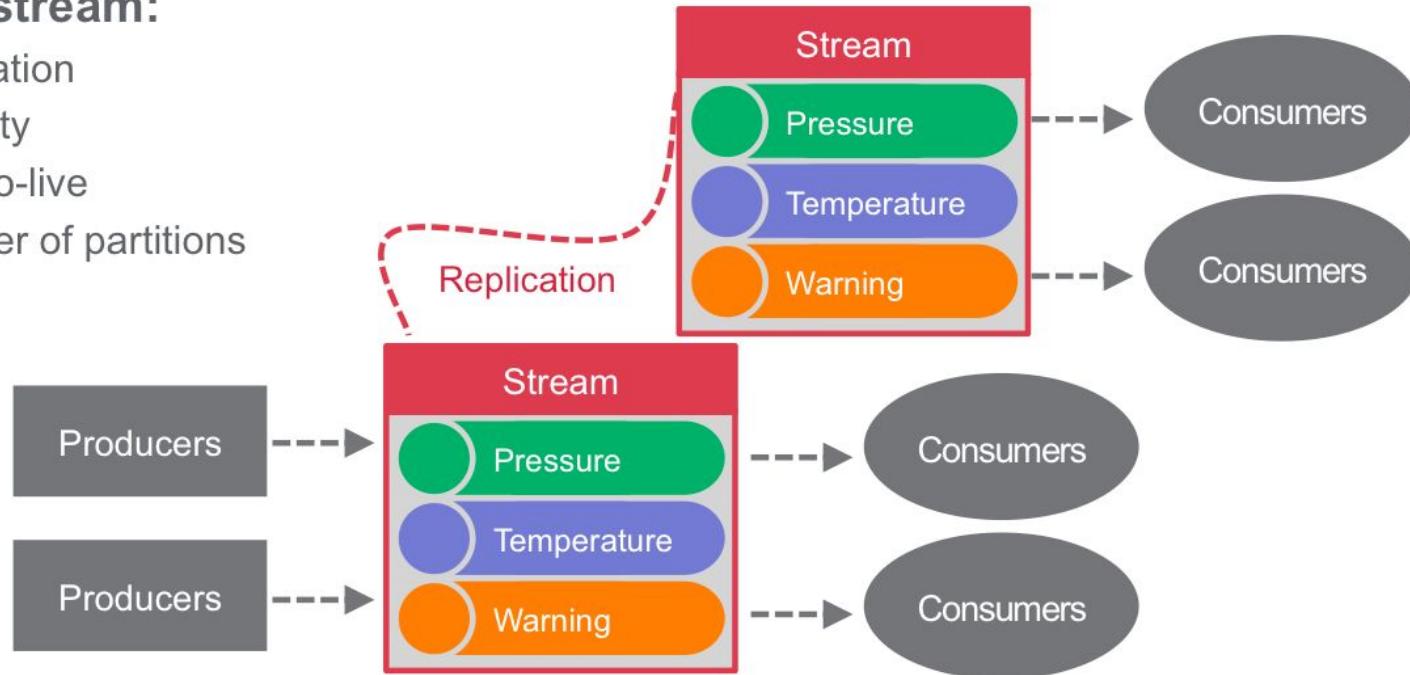


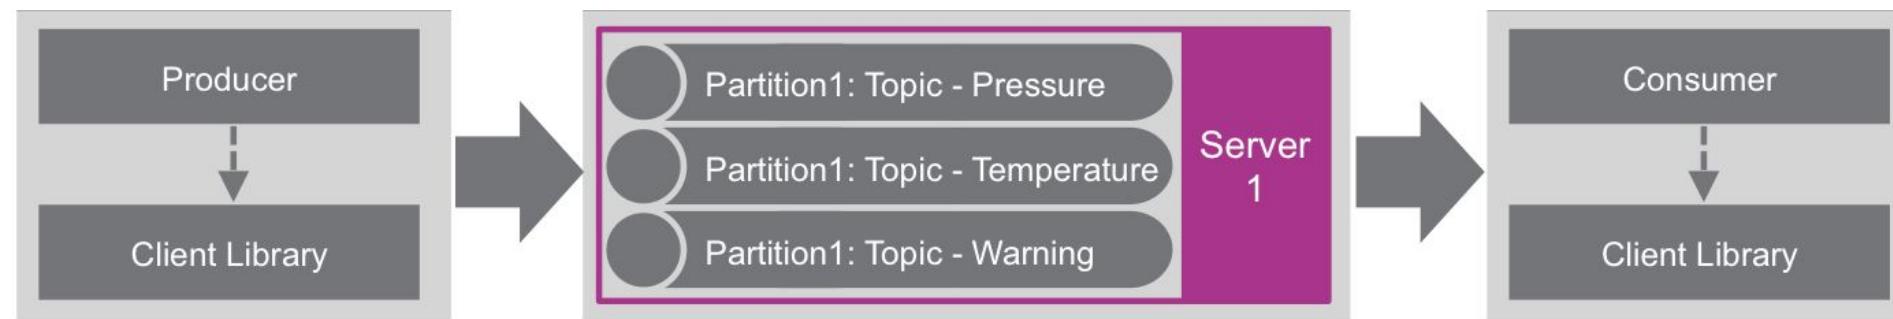
New  
Message

Old  
Message

# Core Components: Streams

- **Stream:**
  - collection of topics managed together
- **Manage stream:**
  - replication
  - security
  - time-to-live
  - number of partitions

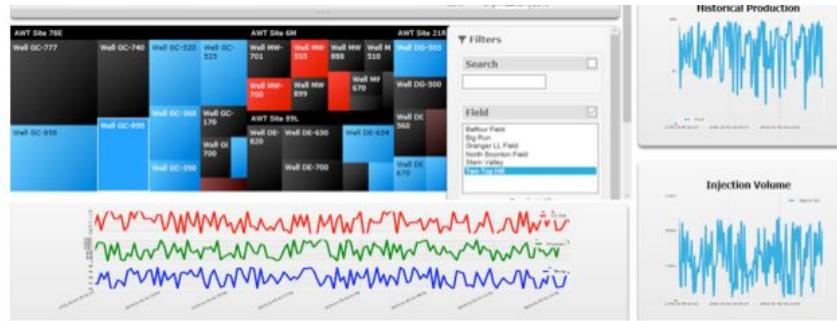
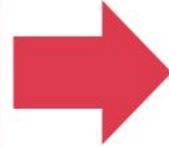
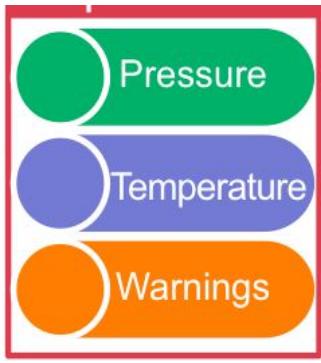
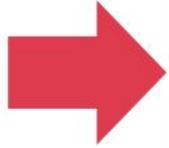


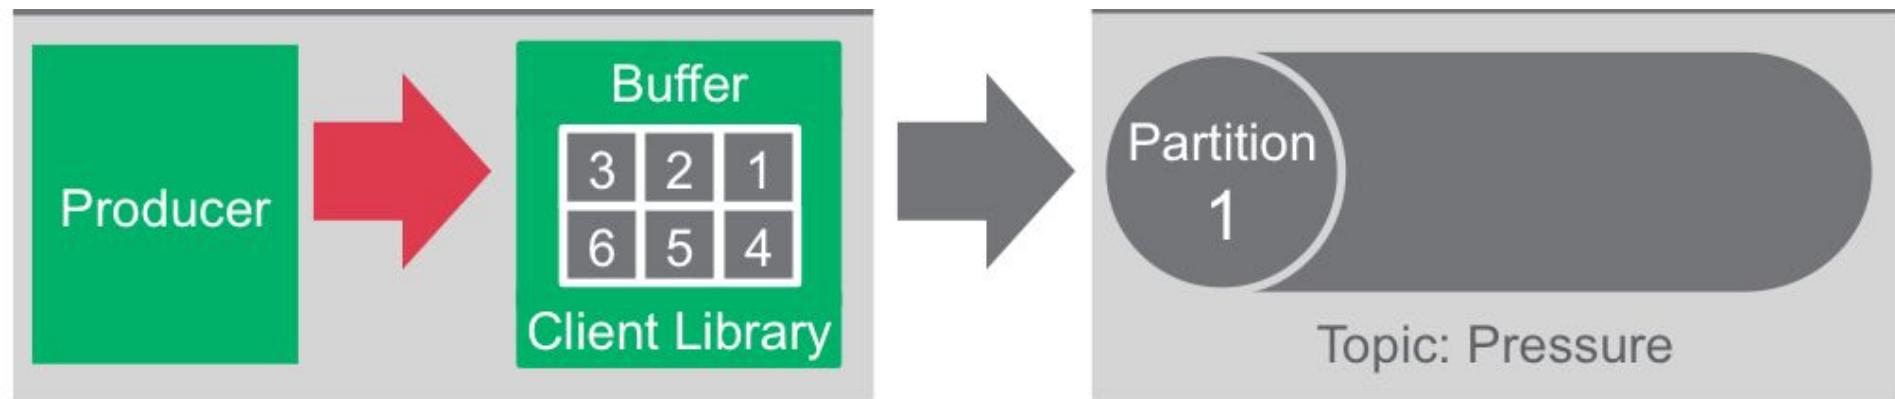


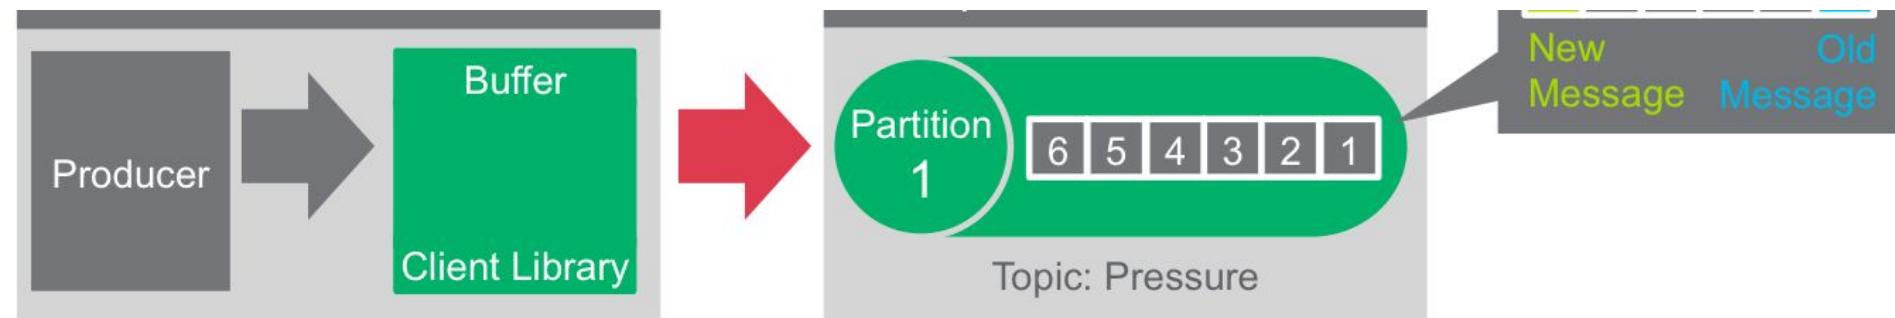


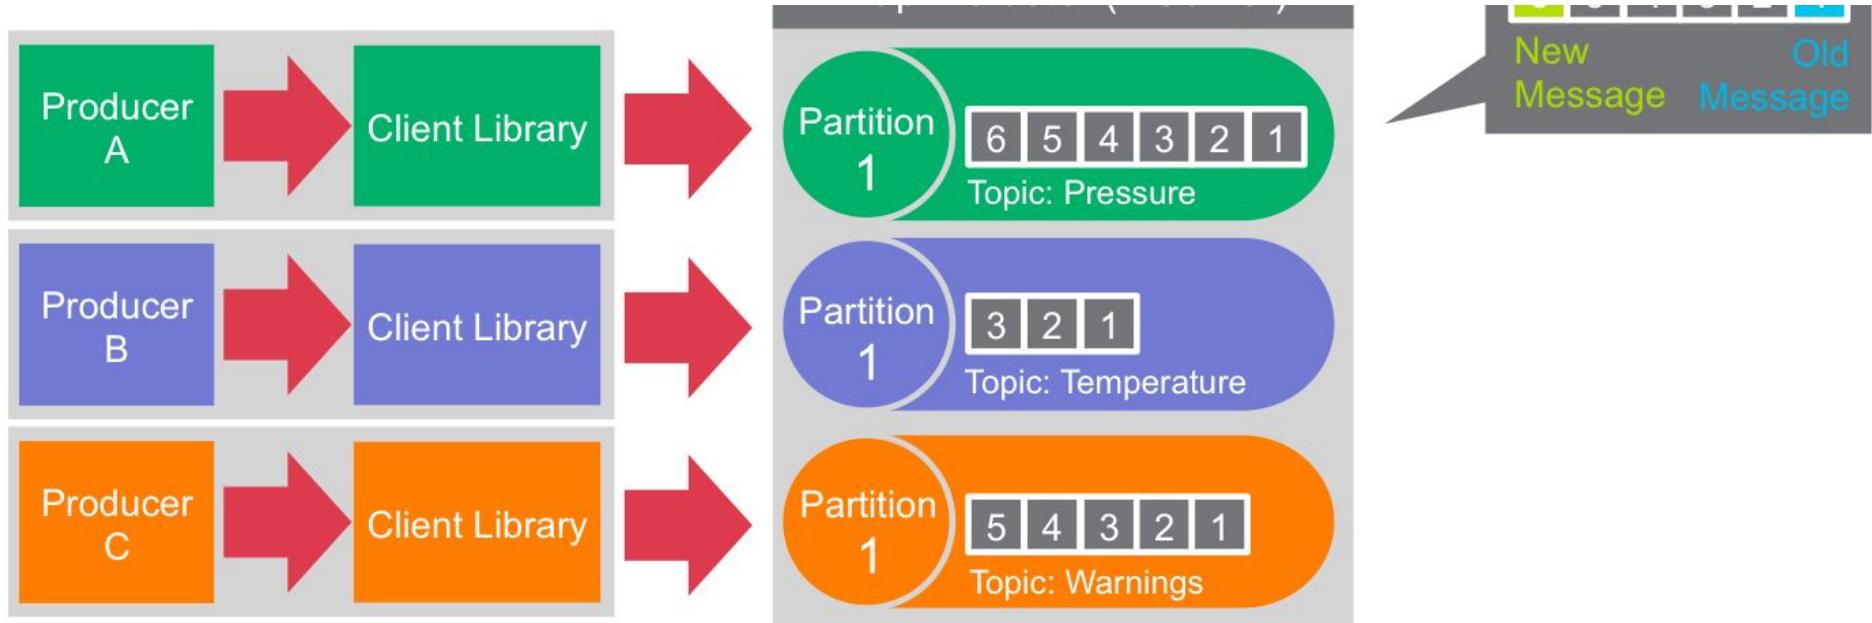


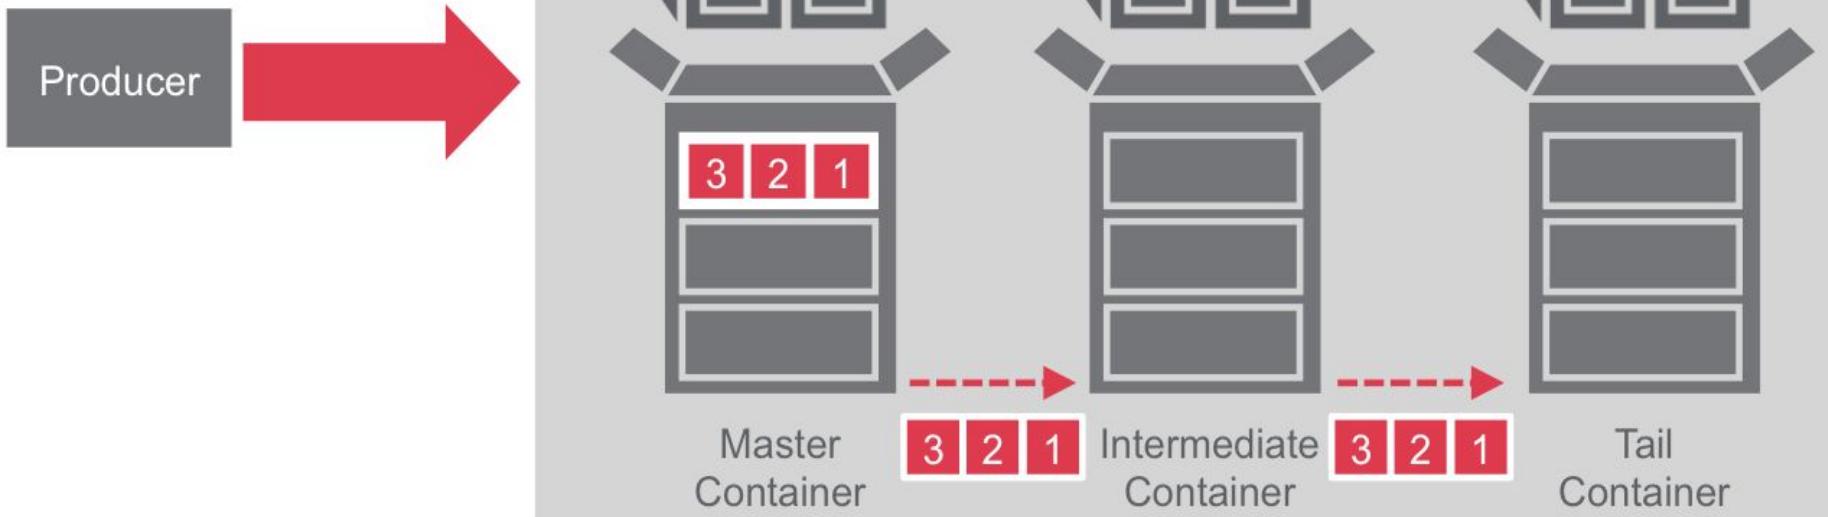
# Goals

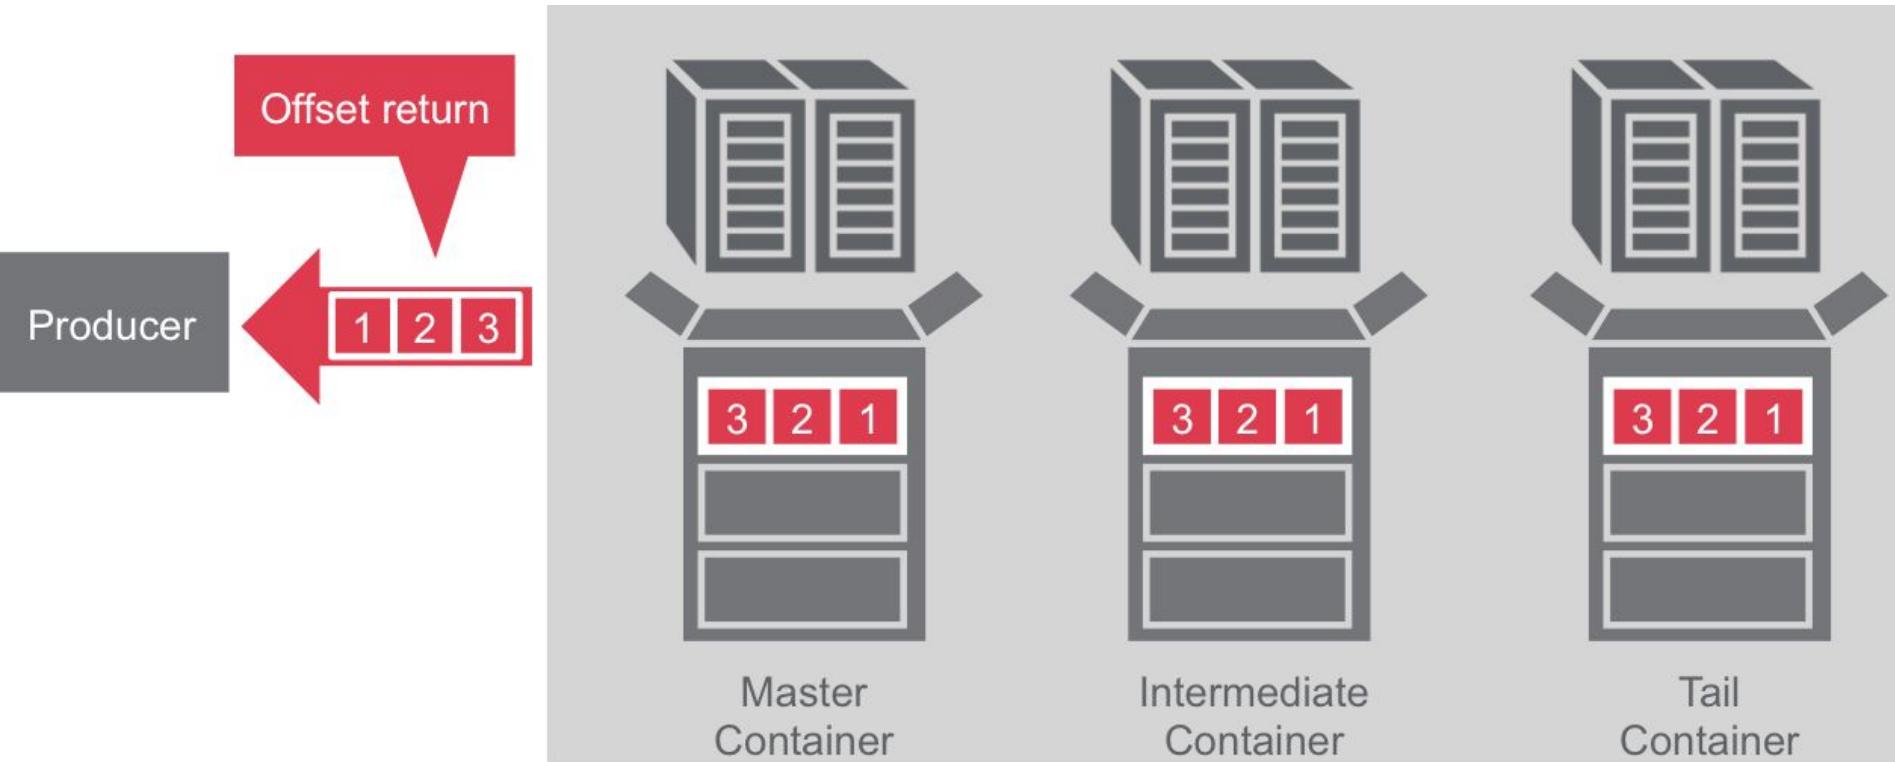




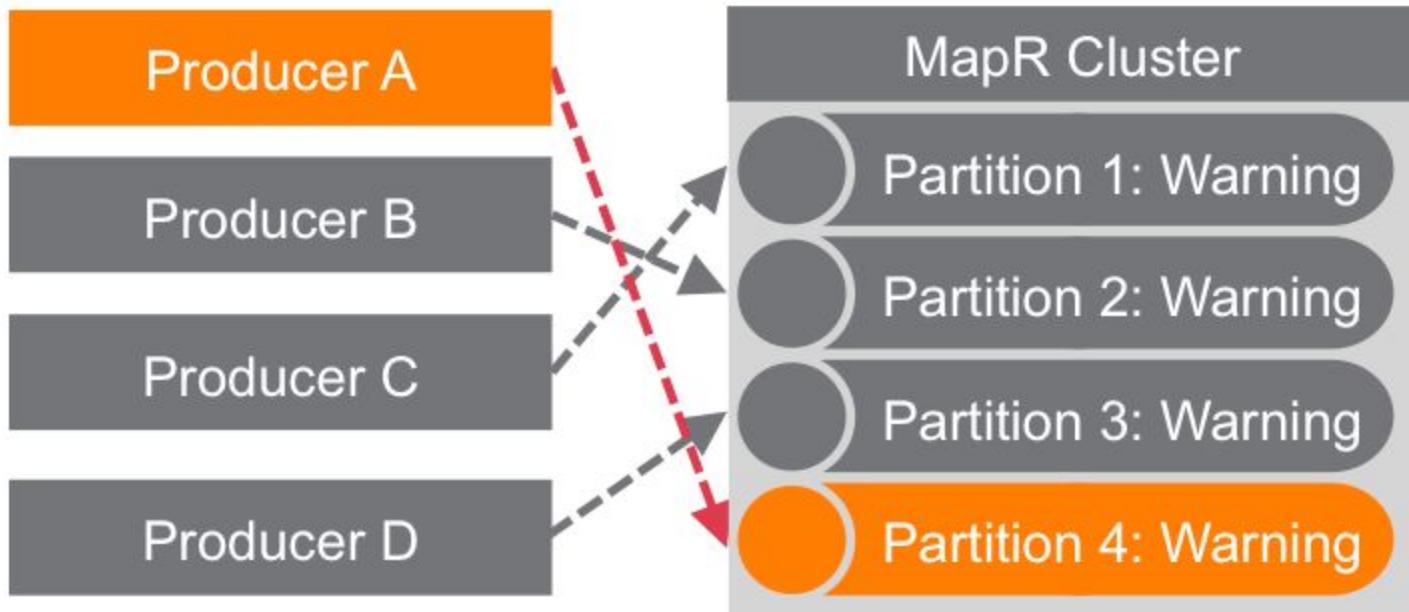








## Example: Producer “A” specifies Partition ID=4

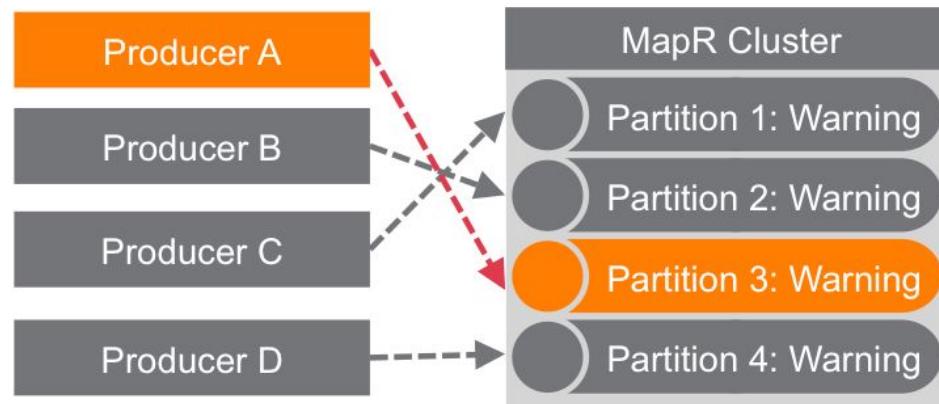


# Parallelism When Publishing

---

How do servers choose which partition to publish to?

**Example: Producer “A” specifies key**

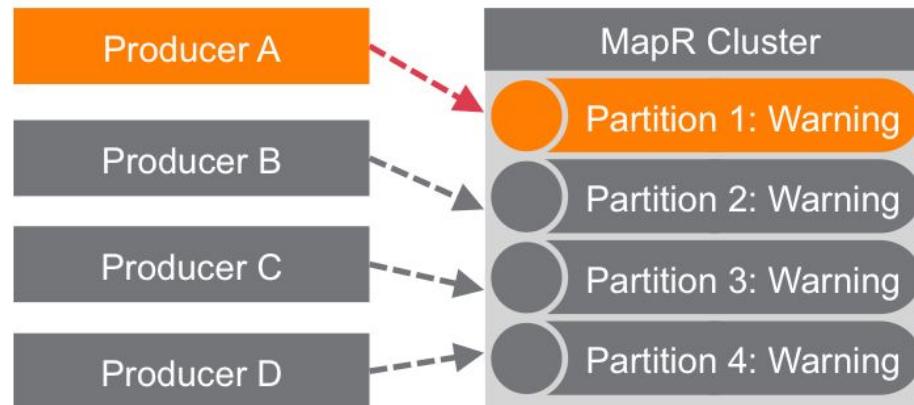


# Parallelism When Publishing

---

How do servers choose which partition to publish to?

**Partitions assigned randomly, sticky round robin**





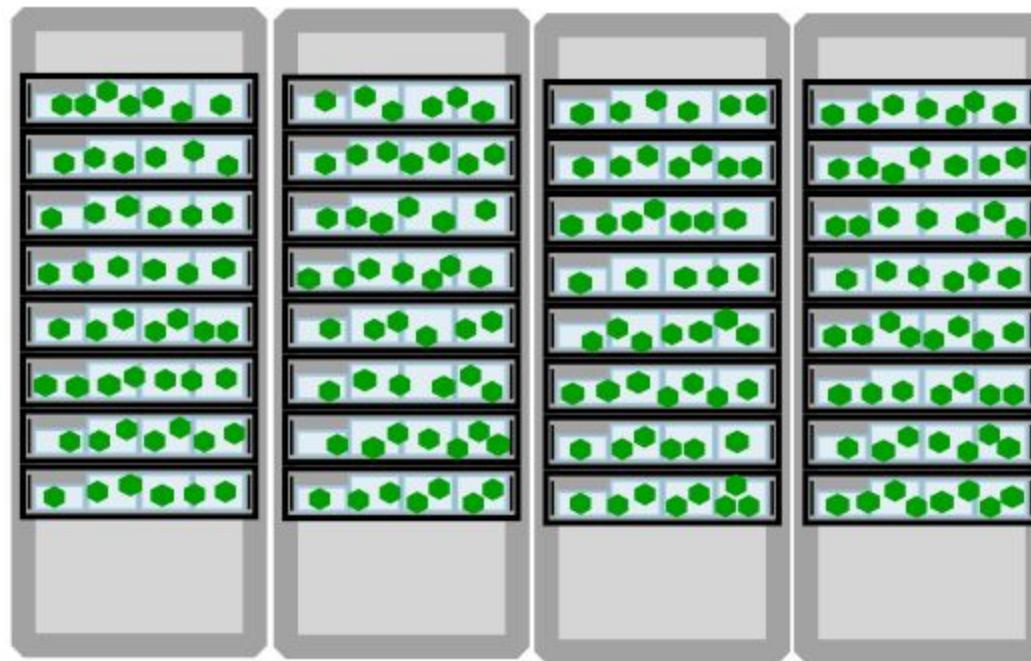


## When Are Messages Deleted?

---

- Older messages deleted based on time to live
- Time-to-live set when stream is created
- Messages don't expire if time-to-live is zero
- Expired messages deleted automatically

## Partitions Are Automatically Distributed



# Parallelism When Reading

---

To read messages in parallel:

- create consumer groups
- consumers with same group.id
- partitions assigned dynamically round-robin

## Partitions Re-Assigned Dynamically

---

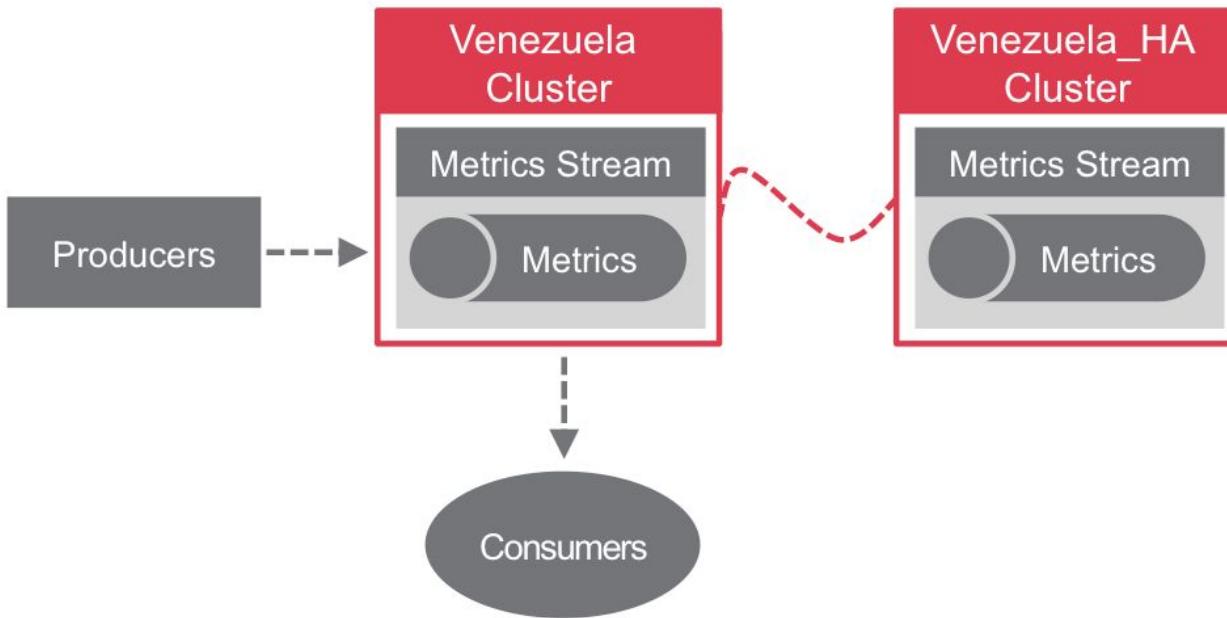
If consumer goes offline, partitions re-assigned

## Saving of Cursor Position: Read Cursor

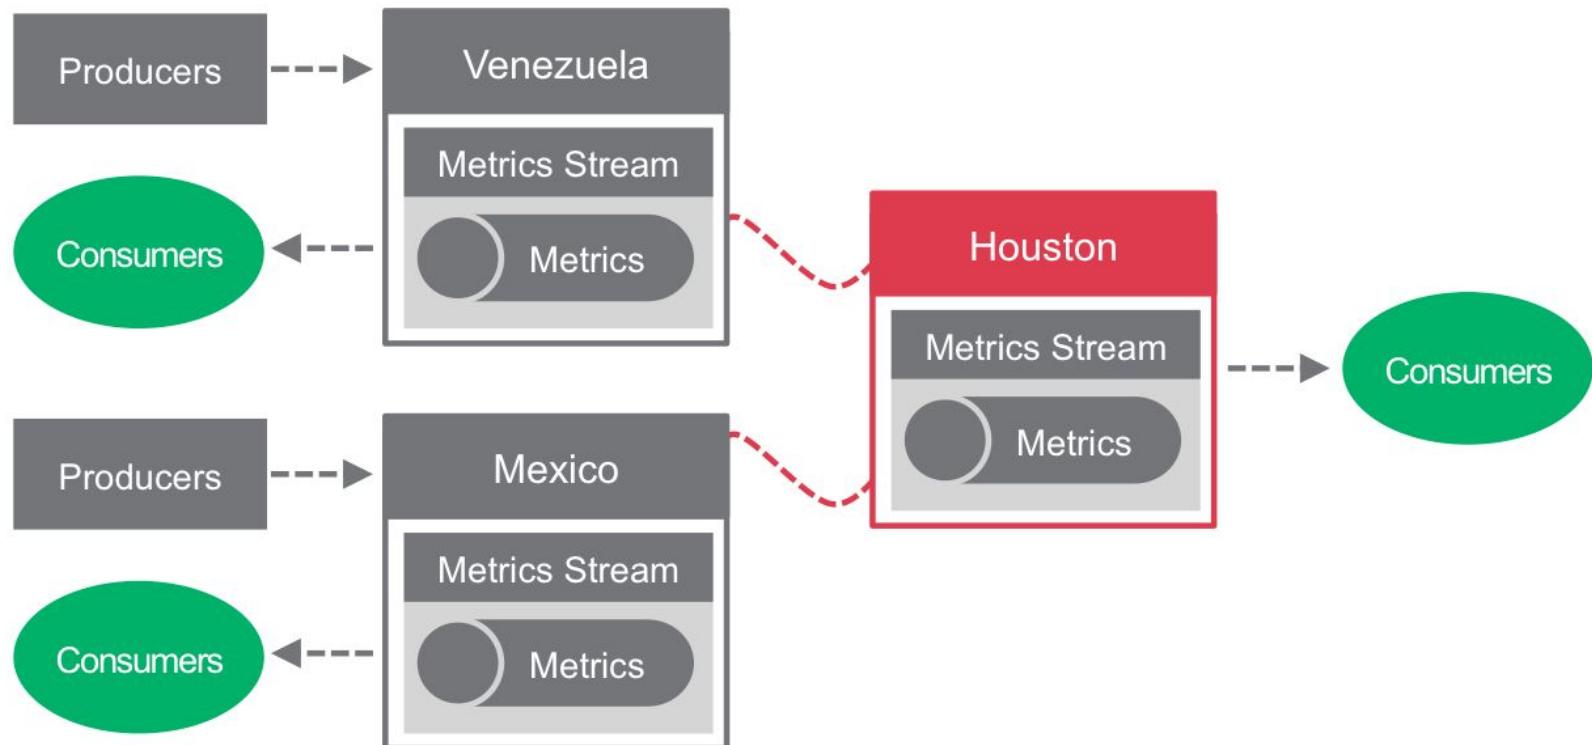
---

- Cursors keep track of messages read
- Read cursor: offset ID of most recent message sent to consumer

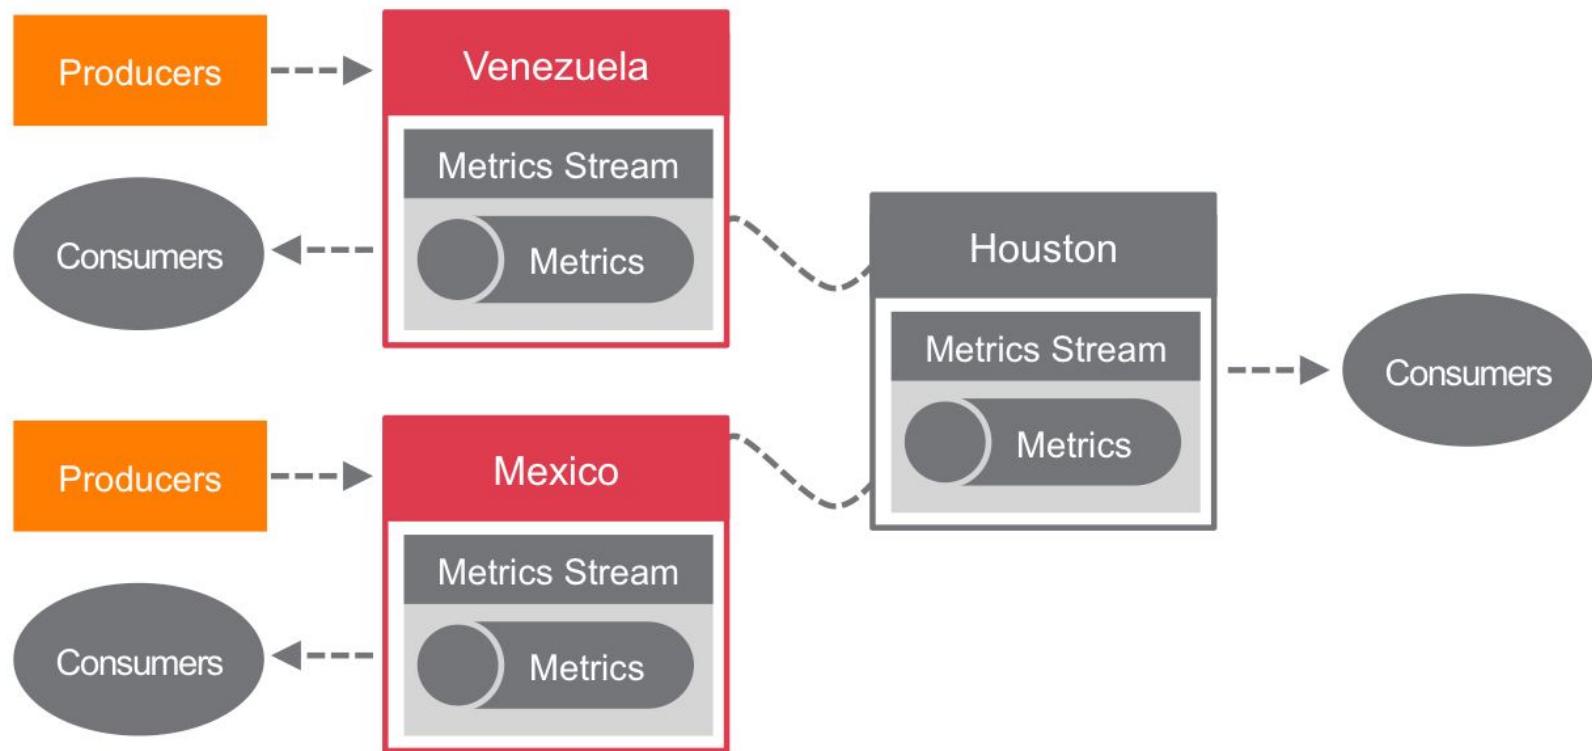
# Replicating Streams: Master-Slave Replication



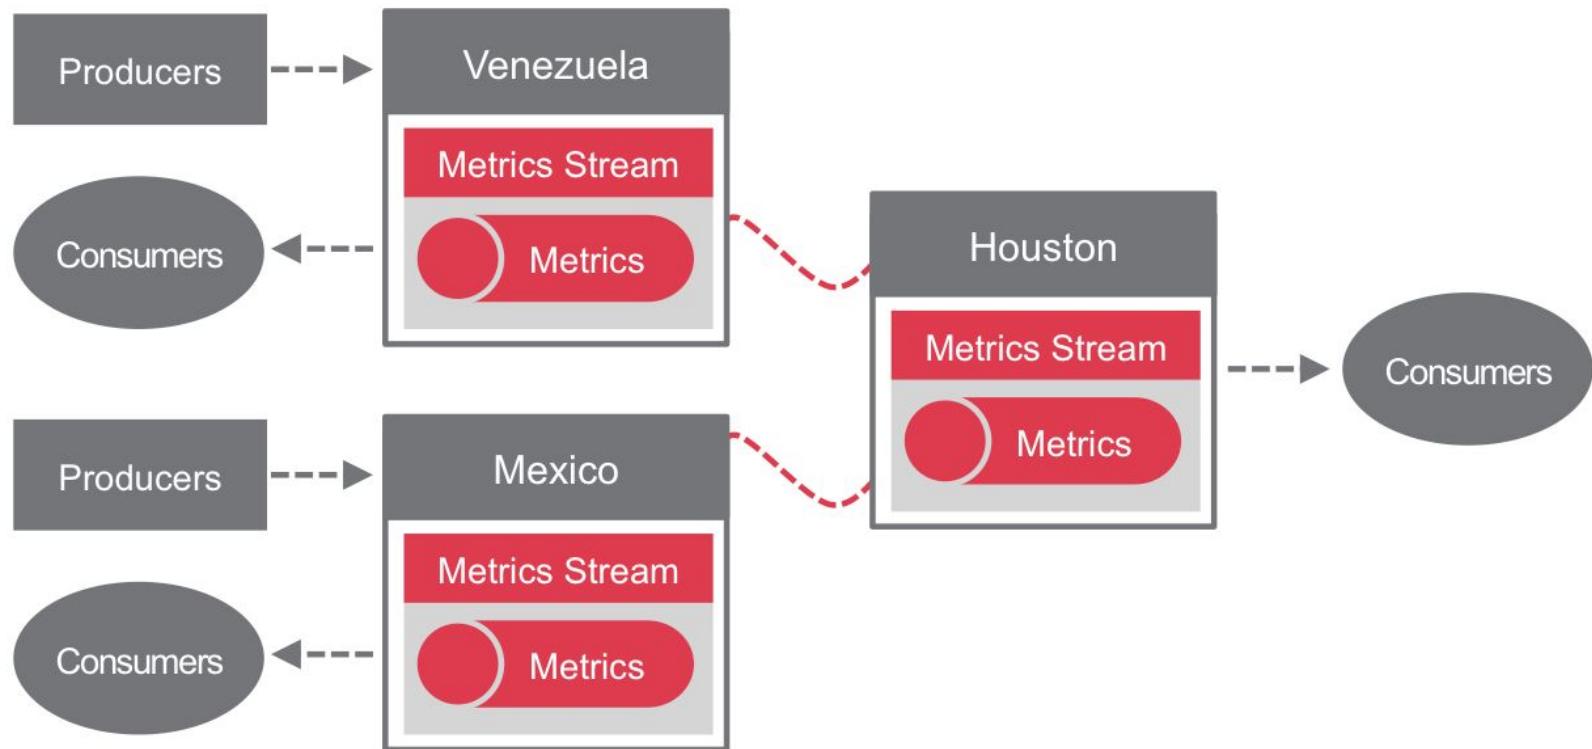
# Replicating Streams: Many-to-One Replication



# Replicating Streams: Many-to-One Replication

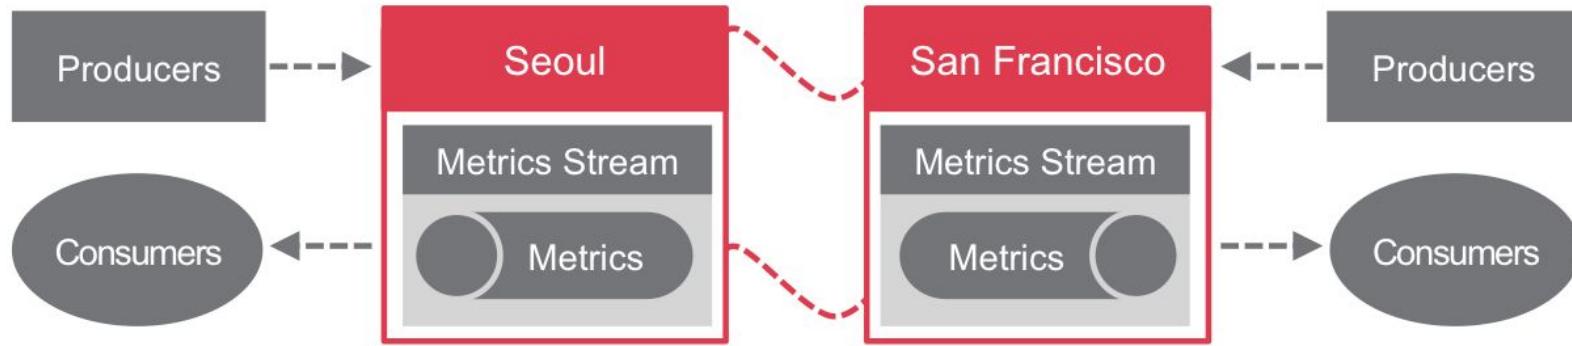


# Replicating Streams: Many-to-One Replication



# Replicating Streams: Multi-Master Replication

---

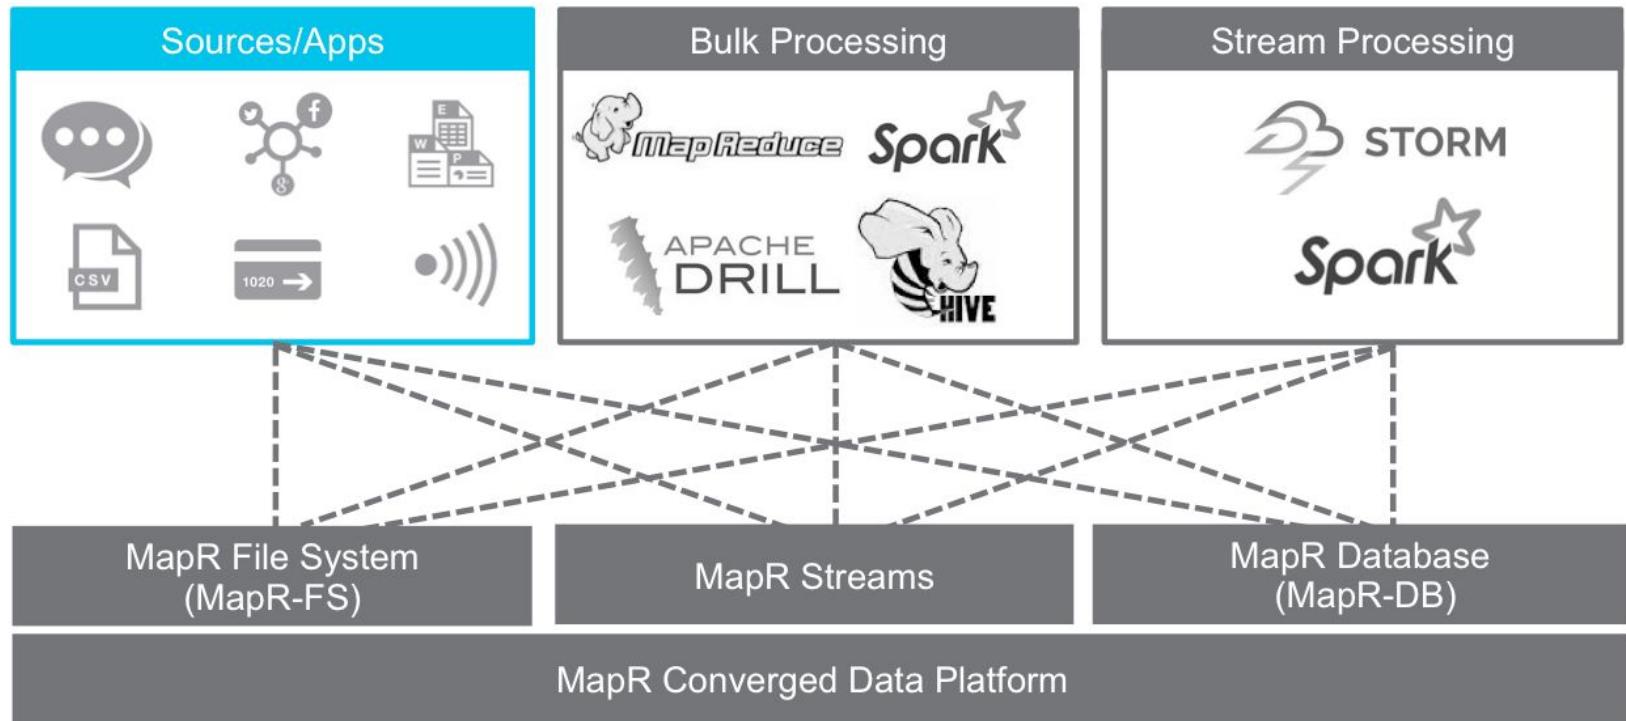


# Stream Replication

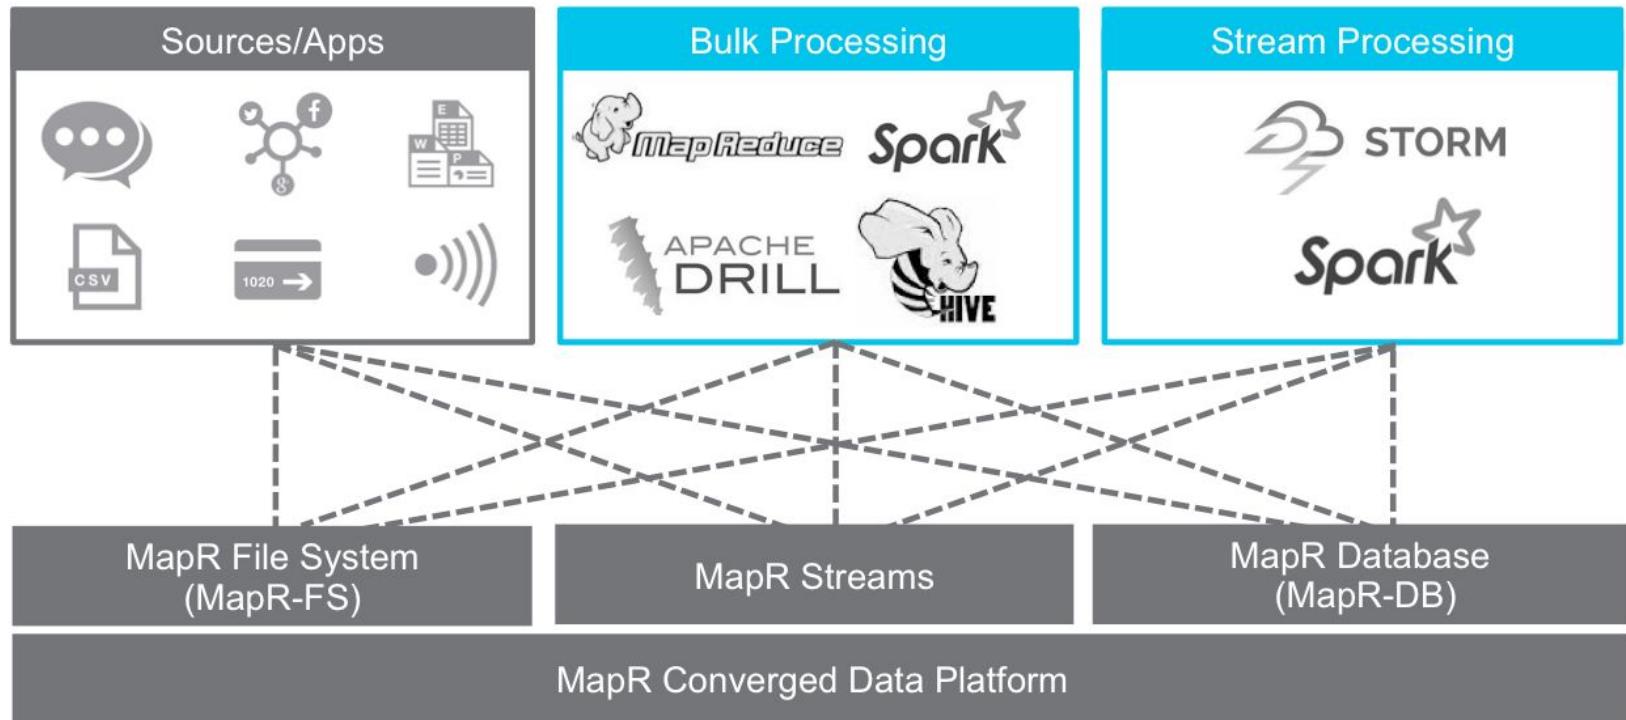
---



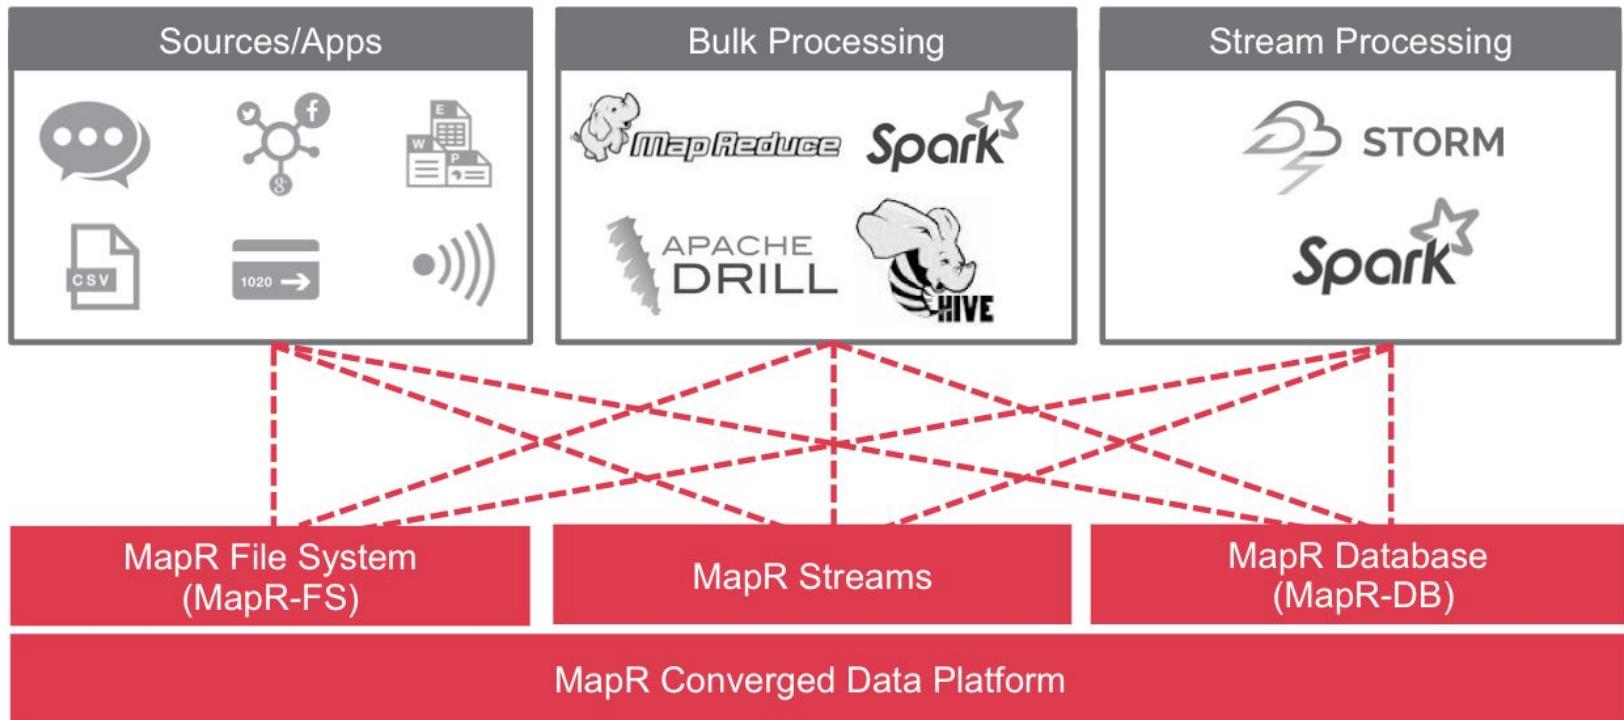
# Building a Complete Data Architecture



# Building a Complete Data Architecture



# Building a Complete Data Architecture



# Goals



# Learning Goals

---

- Create a stream
- Develop a Java producer
- Develop a Java consumer



# Learning Goals

---

- **Create a stream**
- Develop a Java producer
- Develop a Java consumer

## Create a Stream

---

maprcli command for creating a stream

```
maprcli stream create -path <filepath & name> \
    -consumeperm u:<userID> -produceperm u:<userID> \
    -topicperm u:<userID>
```

# Create a Stream

---

Optional parameters to grant security permissions

```
maprcli stream create -path /<filepath & name> \
    -consumeperm u:<userID> -produceperm u:<userID> \
    -topicperm u:<userID>
```

## Create a Stream

---

produceperm determines which users can publish messages to topics in a stream

```
maprcli stream create -path /<filepath & name> \
    -consumeperm u:<userID> -produceperm u:<userID> \
    -topicperm u:<userID>
```

## Create a Stream

---

consumeperm determines which users can read topics from a stream

```
maprcli stream create -path /<filepath & name> \
    -consumeperm u:<userID> -produceperm u:<userID> \
    -topicperm u:<userID>
```

## Create a Stream

---

topicperm determines which users can create/remove topics

```
maprcli stream create -path /<filepath & name> \
    -consumeperm u:<userID> -produceperm u:<userID> \
    -topicperm u:<userID>
```

## Create a Stream

---

defaultpartitions determines how many partitions are created in a topic

```
maprcli stream create -path /<filepath & name> \
    -consumeperm u:<userID> -produceperm u:<userID> \
    -defaultpartitions 3
```

# Optionally Create a Topic

---

## Create topics manually

```
maprcli stream topic create \
    -path <path and name of the stream> \
    -topic <name of the topic>
```

# Goals



# Learning Goals

---

- Create a stream
- **Develop a Java producer**
- Develop a Java consumer

# Producer Sending Messages

---

1. Set producer properties

2. Create a producer

3. Build the message

4. Send the message

## Sample Producer: Set Producer Properties

```
// Set up KafkaProducer properties

Properties properties = new Properties();

properties.put("key.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");

properties.put("value.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");

. . .
```

# Some Producer Properties

---

Property	Description
key.serializer	Java class to use to serialize the key of each message
value.serializer	Java class to use to serialize the value of each message
client.id	ID that identifies the producer (optional)

# Producer Sending Messages

---

1. Set producer properties

2. **Create a producer**

3. Build the message

4. Send the message

## Sample Producer: Create a Producer

---

```
import org.apache.kafka.clients.producer.KafkaProducer;  
  
public static KafkaProducer producer;  
  
// Instantiate KafkaProducer with properties  
producer = new KafkaProducer<String, String>(properties);
```

# Producer Sending Messages

---

1. Set producer properties
2. Create a producer
- 3. Build the message**
  - Specify:
    - topic – either already exists or is implicitly created
    - value of some type V
  - Optionally specify:
    - key of type K
    - partition ID of type integer
4. Send the message

## Sample Producer: Build Message

---

```
String topic="<streamname>:<topicname>";

for(int i = 0; i < 3; i++) {
    // message text
    String txt = "msg " + i;

    // create ProducerRecord with topic and text
    ProducerRecord<String, String> record =
        new ProducerRecord<String, String>(topic, txt);
    . . .
}
```

# Producer Sending Messages

---

1. Set producer properties
2. Create a producer
3. Build the message
- 4. Send the message**
  - `producer.send(record)`

## Sample Producer: Send the Message

---

```
for(int i = 0; i < 3; i++) {  
    . . .  
    Callback cb = new ProducerCallback();  
    producer.send(record, cb);  
}  
producer.close()
```

## Sample Producer: Send the Message

---

```
for(int i = 0; i < 3; i++) {  
    . . .  
    Callback cb = new ProducerCallback();  
    producer.send(record, cb);  
}  
producer.close()
```

## Sample Producer: Send the Message

---

```
for(int i = 0; i < 3; i++) {  
    . . .  
    Callback cb = new ProducerCallback();  
    producer.send(record, cb);  
}  
producer.close()
```

# Sample Producer: Send the Message

---

## Producer callback

```
class ProducerCallback implements Callback {  
  
    public void onCompletion(RecordMetadata meta, Exception e) {  
        if(e != null)  
            e.printStackTrace();  
        System.out.println("offset is: " +meta.offset());  
    }  
}
```

## Sample Producer: All Together

```
public class SampleProducer {  
    String topic="/streams/pump:warning";  
    public static KafkaProducer producer;  
  
    public static void main(String[] args) {  
        producer=setUpProducer();  
        for(int i = 0; i < 3; i++) {  
            String txt = "msg " + i;  
            ProducerRecord<String, String> rec = new  
                ProducerRecord<String, String>(topic, txt);  
            producer.send(rec);  
            System.out.println("Sent msg number " + i);  
        }  
        producer.close();  
    }  
}
```

# Goals



# Learning Goals

---

- Create a stream
- Develop a Java producer
- **Develop a Java consumer**

# Consumer Getting Messages

---

## 1. Set consumer properties

2. Create a consumer

3. Subscribe to topics

4. Polls from topics

5. Process returned messages

## Sample Consumer: Set Consumer Properties

```
// Set up KafkaConsumer properties
Properties properties = new Properties();
properties.put("key.deserializer",
    "org.apache.kafka.common.serialization.StringDeserializer");
properties.put("value.deserializer",
    "org.apache.kafka.common.serialization.StringDeserializer");
```

## Sample Consumer: Set Consumer Properties

---

```
// Set up KafkaConsumer properties
Properties properties = new Properties();
properties.put("key.deserializer",
    "org.apache.kafka.common.serialization.StringDeserializer");
properties.put("value.deserializer",
    "org.apache.kafka.common.serialization.StringDeserializer");
```

# Some Consumer Properties

Property	Description
key.deserializer value.deserializer	Used to deserialize keys, values
auto.offset.reset	<ul style="list-style-type: none"><li>• <b>Earliest</b>: Reset offset to the offset of earliest message in the partition.</li><li>• <b>Latest (default)</b>: Reset offset to the offset of latest message in the partition.</li><li>• <b>None</b>: Throws NoOffsetForPartitionException exception when consumer next polls for messages in subscription and no offset exists. Consumer must unsubscribe from partition before polling functions correctly</li></ul>

# Consumer Properties

---

`auto.offset.reset:`

`earliest`: reset cursor to the offset of the earliest message in the partition

`latest`: reset the offset to the offset of the latest message in the partition.



# Consumer Getting Messages

---

1. Set consumer properties
2. **Create a consumer**
3. Subscribe to topics
4. Polls from topics
5. Process returned messages

## Sample Consumer: Set up KafkaConsumer

---

```
import org.apache.kafka.clients.consumer.KafkaConsumer;  
  
public static KafkaConsumer consumer;  
  
// Instantiate KafkaConsumer with properties  
consumer = new KafkaConsumer<String, String>(properties);
```

# Consumer Getting Messages

---

1. Set consumer properties
2. Create a consumer
- 3. Subscribe to topics**
4. Polls from topics
5. Process returned messages

## Sample Consumer: Subscribe

---

```
public static String topic = "/stream/pump:warning";  
  
// Subscribe to the topic  
consumer.subscribe(Arrays.asList(topic));
```

# Consumer Getting Messages

---

1. Set consumer properties
2. Create a consumer
3. Subscribe to topics
- 4. Polls from topics**
5. Process returned messages

## Sample Consumer: Poll

---

```
long pollTimeOut = 1000;
while (true) {
    // Request unread messages from the topic
    ConsumerRecords<String, String> msgs =
        consumer.poll(pollTimeOut);
    . . .
}
```

# Consumer Getting Messages

---

1. Set consumer properties
  2. Create a consumer
  3. Subscribe to topics
  4. Polls from topics
- 5. Process returned messages**
- One for each message found in `poll()` `ConsumerRecords<K, V>`
  - Key of type K
  - Value of type V
  - Partition and topic as well

## Sample Consumer: Process Returned Messages

```
ConsumerRecords<String, String> msgs =  
    consumer.poll(pollTimeOut);  
  
// Process returned messages  
Iterator<ConsumerRecord<String, String>> iter =  
    msgs.iterator();
```

## Sample Consumer: Process Returned Messages

```
Iterator<ConsumerRecord<String, String>> iter =  
    msgs.iterator();  
  
// get ConsumerRecord from ConsumerRecord iterator  
while (iter.hasNext()) {  
    ConsumerRecord<String, String> record = iter.next();  
    System.out.println("read msg" + record.toString());  
}
```

Example record.toString()

```
topic=/events:sensor,partition=0,offset=192,  
key=null,value=Msg1
```

# Sample Consumer: All Together

```
public class MyConsumer {  
    public static String topic = "/stream/pump:warning";  
    public static KafkaConsumer consumer;  
    public static void main(String[] args) {  
        configureConsumer(args);  
        consumer.subscribe(topic);  
        while (true) {  
            ConsumerRecords<String, String> msg=  
                consumer.poll(pollTimeOut);  
            Iterator<ConsumerRecord<String, String>> iter =  
                msg.iterator();  
            while (iter.hasNext()) {  
                ConsumerRecord<String, String> record = iter.next();  
                System.out.println("read " + record.toString());  
            }  
        }  
        consumer.close();  
    }  
}
```

## Sample Consumer: Run the Consumer

---

```
$java -cp my-app.jar:`mapr classpath` solution.MyConsumer  
  
read msg(topic=/  
events:sensor,partition=0,offset=192,key=null,value=Msg1)  
read msg(topic=/  
events:sensor,partition=0,offset=193,key=null,value=Msg2)  
read msg(topic=/  
events:sensor,partition=0,offset=194,key=null,value=Msg3)
```

# Goals



# Learning Goals

---

- Describe producer properties and options for:
  - buffering and batching of messages
  - publishing to partitions
- Describe consumer properties and options for:
  - fetching data
  - consumer groups
  - read cursors
- Explain messaging semantics



# Learning Goals

---

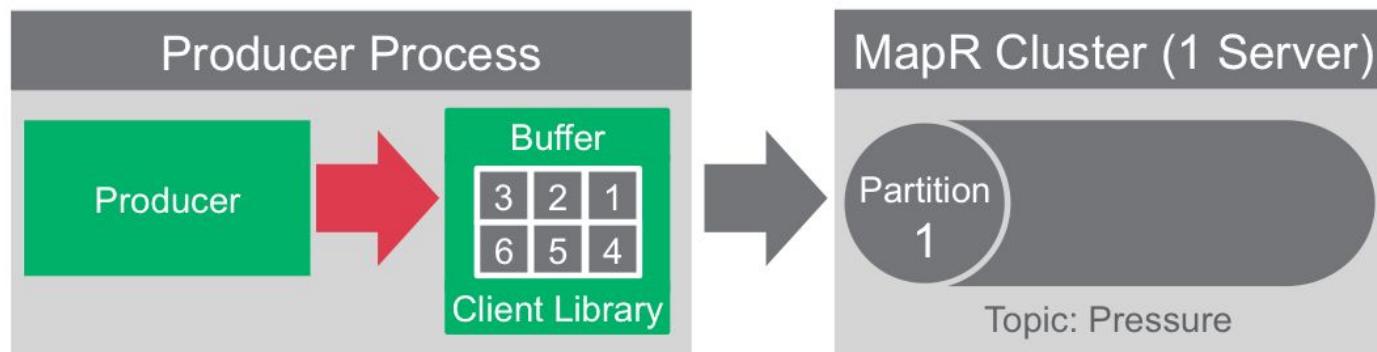
- **Describe producer properties and options for:**
  - buffering and batching of messages
  - publishing to partitions
- Describe consumer properties and options for:
  - fetching data
  - consumer groups
  - read cursors
- Explain messaging semantics

# Review



## Review: How Are Messages Sent?

- Producers create and send messages to client library
- Client library buffers the messages



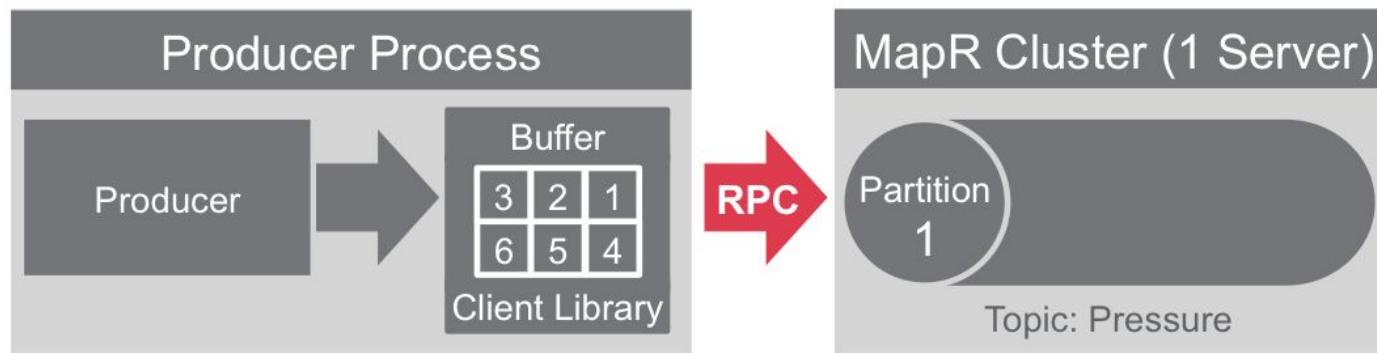
## Producer Buffering and Batching Properties

Property	Description
<b>buffer.memory</b> (supported from Apache Kafka)	(default: 33554432) The total bytes of memory the producer can use to buffer records waiting to be sent to the server. If records are generated faster than they can be delivered to the server the producer will block.

# When Are Messages Published?

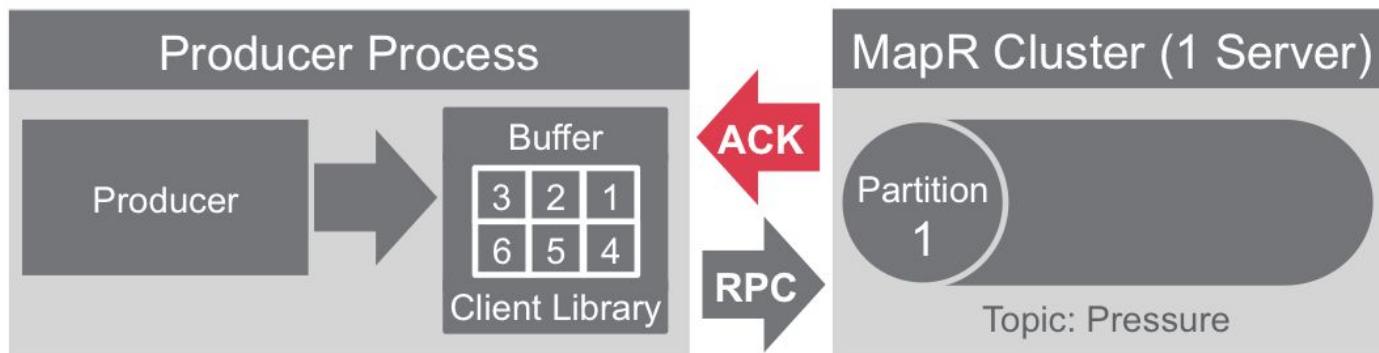
When any of four conditions met:

- Messages buffered is enough to make efficient RPC to Server
- Messages buffered = streams.buffer.max.time.ms setting
- Bytes buffered = buffer.memory setting
- Producer calls flush() method



# Two Modes of Publishing

- streams.parallel.flushers.per.partition **default true:**
  - does **not** wait for ACK before sending more messages
  - possible for messages to arrive out of order
- streams.parallel.flushers.per.partition **set to false:**
  - client library will wait for ACK from server
  - slower than default setting



# Guaranteed Ordering

---

- Messages are stored in partitions in the order they are received
- Messages are always read from a partition in order

## Note:

- By default producer buffers are flushed in parallel
- Can cause out of order messages
- If an issue, MapR Streams:  
`streams.parallel.flushers.per.partition=false`



# Goals



# Learning Goals

---

- **Describe producer properties and options for:**
  - buffering and batching of messages
  - **publishing to partitions**
- Describe consumer properties and options for:
  - fetching data
  - consumer groups
  - read cursors
- Explain messaging semantics

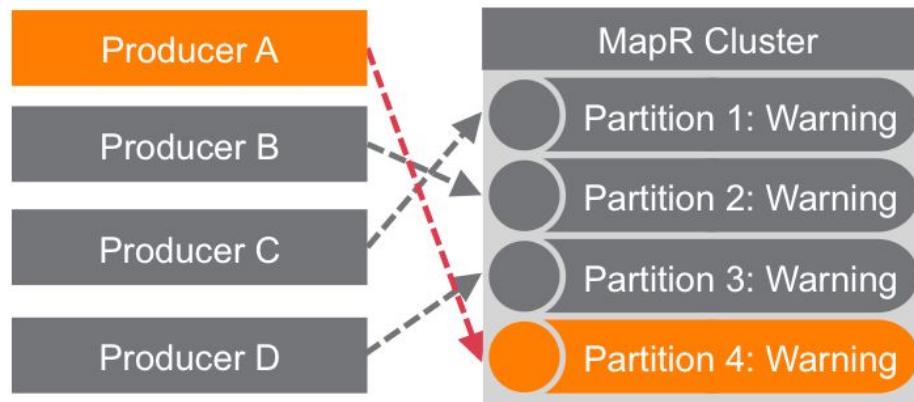
# Review



# Review: Parallelism When Publishing

How do servers choose which partition to publish to?

**Example: Producer “A” specifies Partition ID=4**

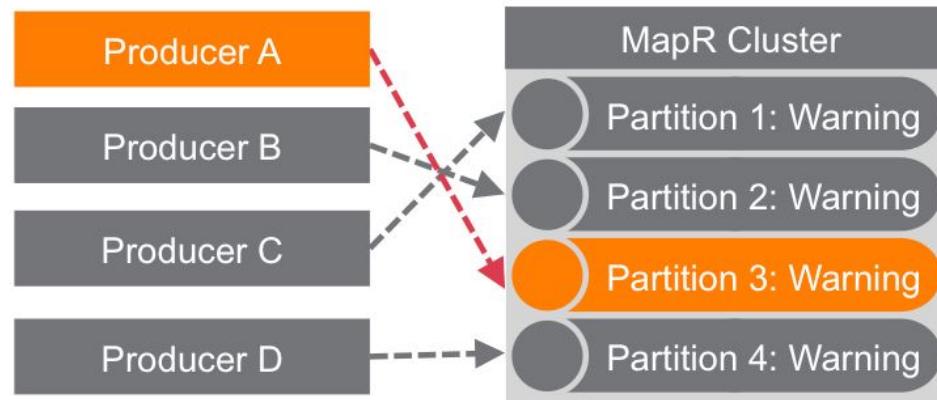




# Review: Parallelism When Publishing

How do servers choose which partition to publish to?

## Example: Producer “A” specifies key

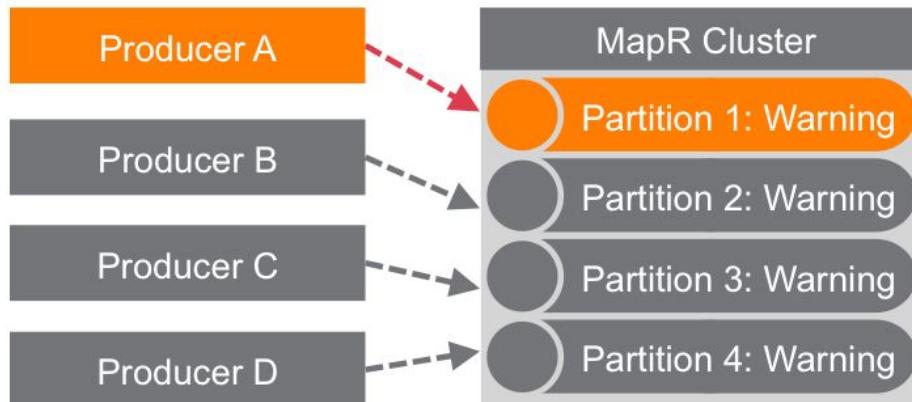




# Review: Parallelism When Publishing

How do servers choose which partition to publish to?

**Partitions assigned randomly, sticky round robin**



## Producer Builds Message Specifying Key for Partitioning

---

```
String key="sensorxyz";
// create ProducerRecord with topic and text
ProducerRecord<String, String> record =
    new ProducerRecord<String, String>(topic, key, txt);

producer.send(record);
```

## Producer Builds Message Specifying a Partition

---

```
Integer partition=1;  
// create ProducerRecord with topic and text  
ProducerRecord<String, String> record =  
    new ProducerRecord<String, String>(topic, partition, txt);  
  
producer.send(record);
```

## Producer Property for Topic Metadata

Property	Description
<code>metadata.max.age.ms</code> (supported from Apache Kafka)	How frequently to fetch metadata, which might discover new topics or partitions (default 600 * 1000 msec)

## Producer Property for Streams Partitioner Class

Property	Description
<code>streams.partitionner.class</code>	Allows to specify a java class for determining topic partition

```
Properties properties = new Properties();
properties.put("streams.partitionner.class",
"myorg.MyPartitioner.class");
```

## Example: Partitioner Implementation

```
public class MyPartitioner implements Partitioner {  
    public int partition(String topic, Object key,  
                        byte[] keyBytes, Object value,  
                        byte[] valueBytes, Cluster cluster) {  
        List topicPartitions=cluster.partitionsForTopic(topic);  
        int partition;  
        partition=key % topicPartitions.size()  
  
        return partition;  
    }  
}
```

# Goals



# Learning Goals

---

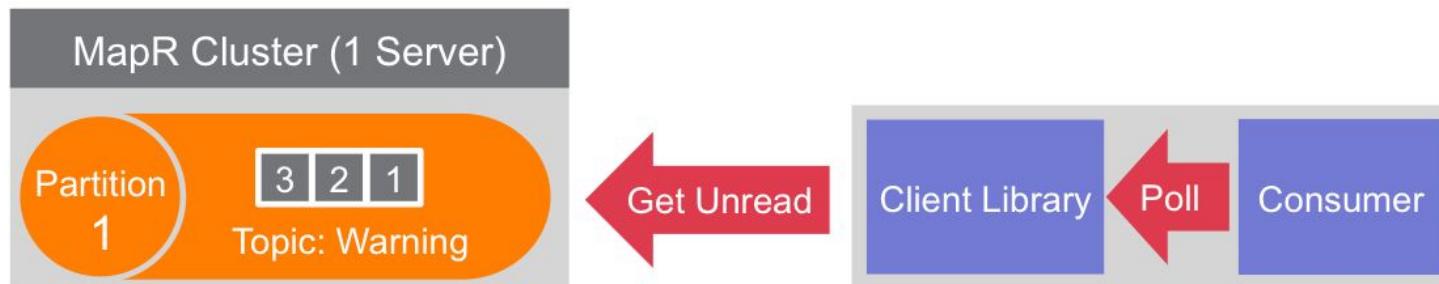
- Describe producer properties and options for:
  - buffering and batching of messages
  - publishing to partitions
- **Describe consumer properties and options for:**
  - fetching data
  - consumer groups
  - read cursors
- Explain messaging semantics

# Review



# Review: How are Messages Read?

- Consumers read at their own pace
- Consumers polls topics it subscribes to
- Client library requests unread messages



## Some Consumer Properties

---

Property	Description
<code>max.partition.fetch.bytes</code>	Number of bytes of message data to attempt to fetch for each partition in each poll request (default: 64KB)
<code>fetch.min.bytes</code>	The minimum amount of data the server should return for a fetch request

# Goals



# Learning Goals

---

- Describe producer properties and options for:
  - buffering and batching of messages
  - publishing to partitions
- **Describe consumer properties and options for:**
  - fetching data
  - **consumer groups**
  - read cursors
- Explain messaging semantics

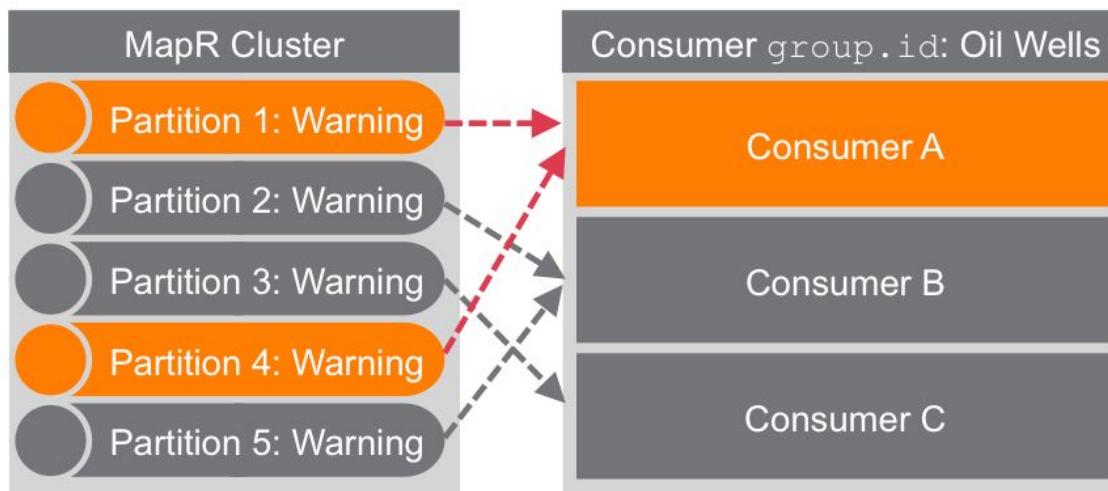
# Review



# Review: Parallelism When Reading

To read messages in parallel:

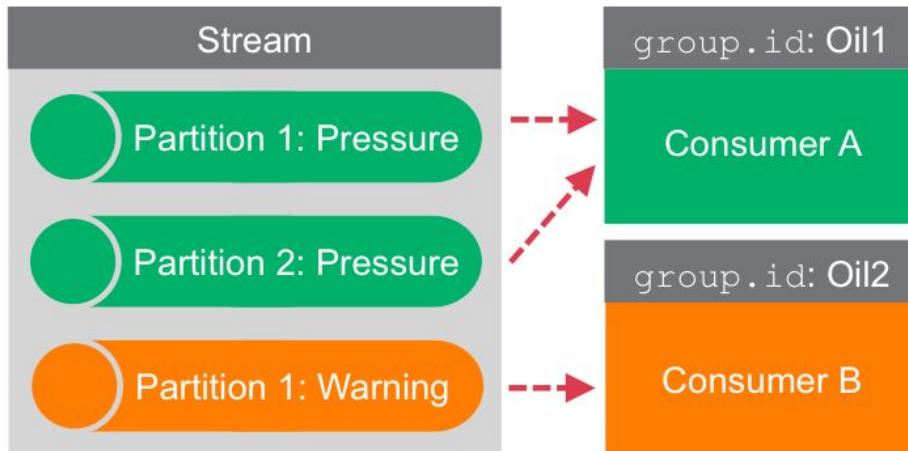
- create consumer groups
- consumers with same `group.id`
- partitions assigned dynamically round-robin



# Consumer Groups

---

Enable cursor persistence which is useful even if group has only one member



# Consumer Group Property

Property	Description
group.id	identifies the group of consumer processes to which this consumer belongs

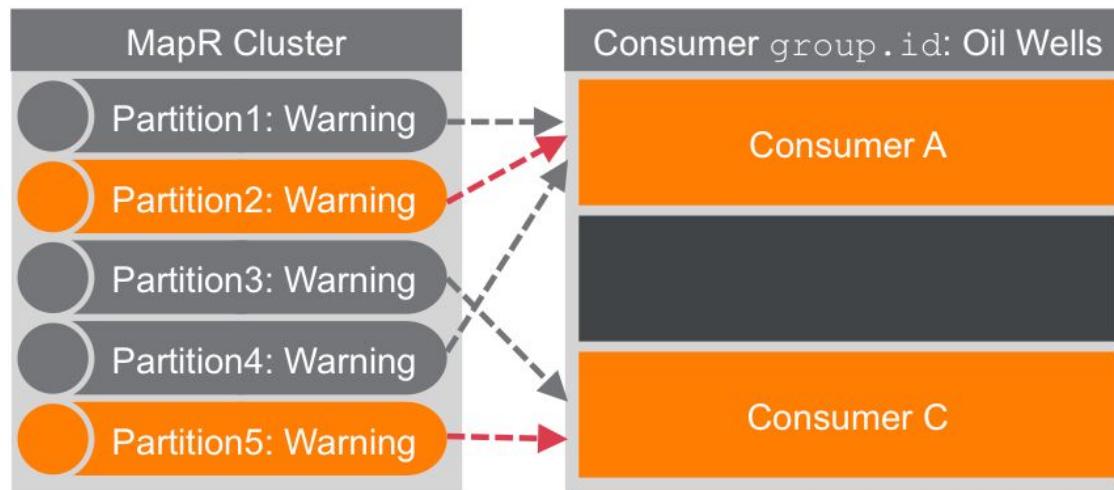
```
// Set up KafkaConsumer properties
Properties props = new Properties();
props.put("group.id", "pumppressure");
consumer = new KafkaConsumer<String, String>(props);
```

# Review



## Review: Partitions Re-assigned Dynamically

If consumer goes offline, partitions re-assigned



## Example: ConsumerRebalanceListener

```
public class Listener implements ConsumerRebalanceListener {  
    private Consumer<?, ?> consumer;  
    // Constructor  
    public Listener(Consumer<?, ?> consumer) {  
        this.consumer = consumer;  
    }  
    public void onPartitionsAssigned(  
        Collection<TopicPartition> partitions) {  
    }  
    public void onPartitionsRevoked(  
        Collection<TopicPartition> partitions) {  
    }  
}
```

## Example: ConsumerRebalanceListener

```
public class Listener implements ConsumerRebalanceListener {  
    private Consumer<?, ?> consumer;  
    // Constructor  
    public Listener(Consumer<?, ?> consumer) {  
        this.consumer = consumer;  
    }  
    public void onPartitionsAssigned(  
        Collection<TopicPartition> partitions) {  
    }  
    public void onPartitionsRevoked(  
        Collection<TopicPartition> partitions) {  
        this.consumer.commitSync();  
    }  
}
```

## Example: ConsumerRebalanceListener

```
public class Listener implements ConsumerRebalanceListener {  
    private Consumer<?, ?> consumer;  
  
    public void onPartitionsAssigned(  
        Collection<TopicPartition> partitions) {  
    }  
    public void onPartitionsRevoked(  
        Collection<TopicPartition> partitions) {  
        this.consumer.commitSync();  
        System.out.print("Stopped Listening to:" + partitions);  
    }  
}
```

## Example: ConsumerRebalanceListener

```
public class Listener implements ConsumerRebalanceListener {  
    private Consumer<?, ?> consumer;  
  
    public void onPartitionsAssigned(  
        Collection<TopicPartition> partitions) {  
        System.out.print("Started Listening to:" + partitions);  
    }  
    public void onPartitionsRevoked(  
        Collection<TopicPartition> partitions) {  
    }  
}
```

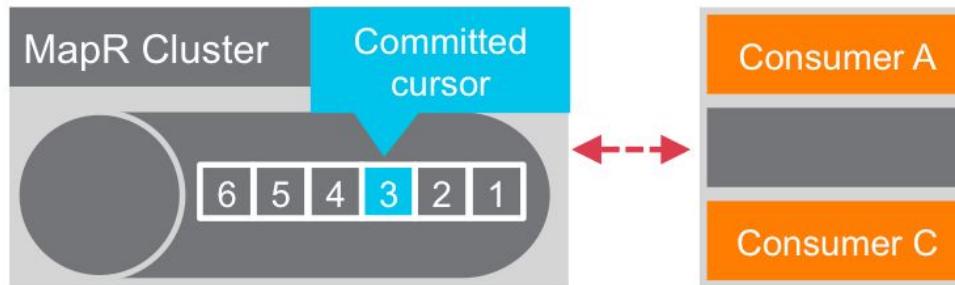
## Example: ConsumerRebalanceListener

```
// subscribe to topic, provide ConsumerRebalanceListener
consumer.subscribe(topic, new Listener(consumer));
```

# Partition Balancing and the Committed Cursor

When new consumers join, leave or new partitions added, rebalance occurs

- there may be a slight pause in message delivery while this happens
- can cause duplicates if consumers don't properly commit cursors



# Goals



# Learning Goals

---

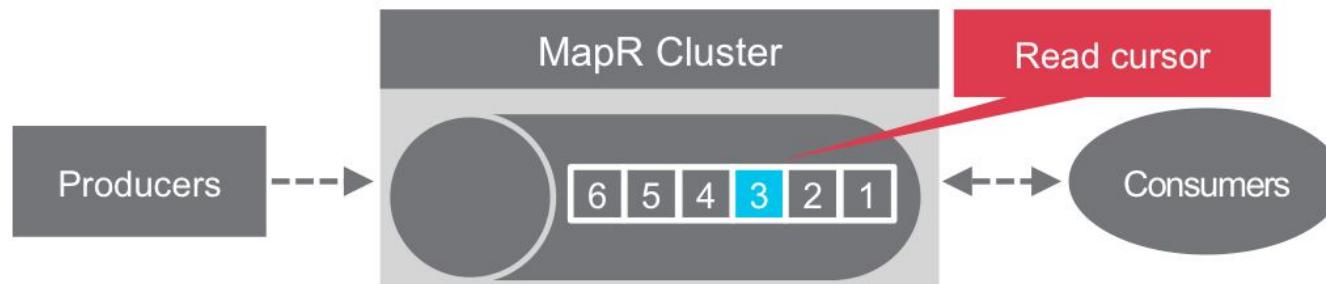
- Describe producer properties and options for:
  - buffering and batching of messages
  - publishing to partitions
- **Describe consumer properties and options for:**
  - fetching data
  - consumer groups
  - **read cursors**
- Explain messaging semantics

# Review



## Review: Saving of Read Cursor Position

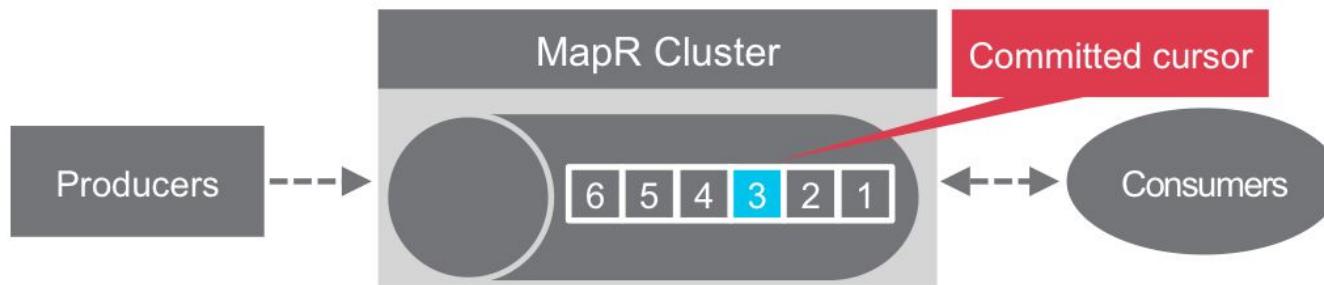
- Cursors keep track of messages read
- Read cursor: offset ID of most recent message sent to consumer





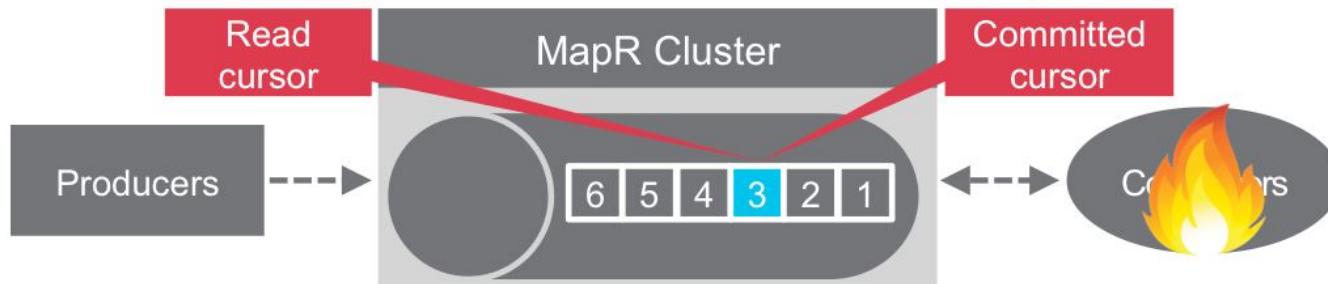
# Review: Saving of Committed Cursor Position

Committed cursor: saved current position of read cursor



# How Often Should Consumer Commit Cursor?

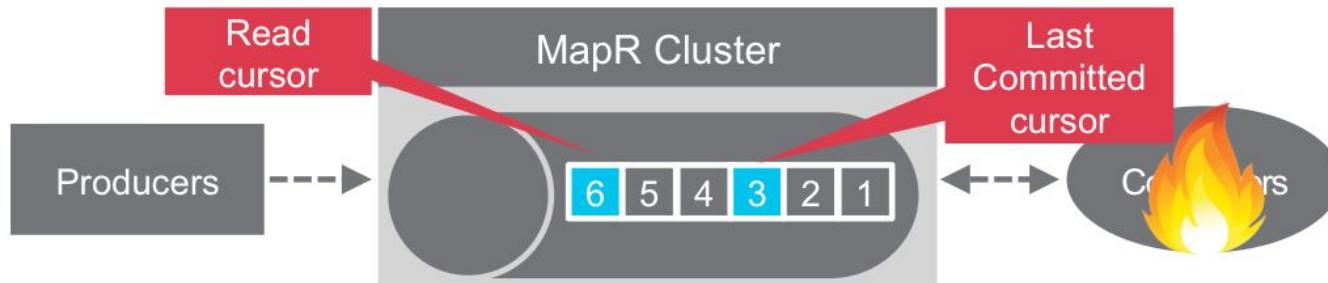
More frequent commits = less read duplication



# Read duplication when Consumer fails?

How much read duplication when consumer fails:

- Time since last consumer commit
- Rate messages published

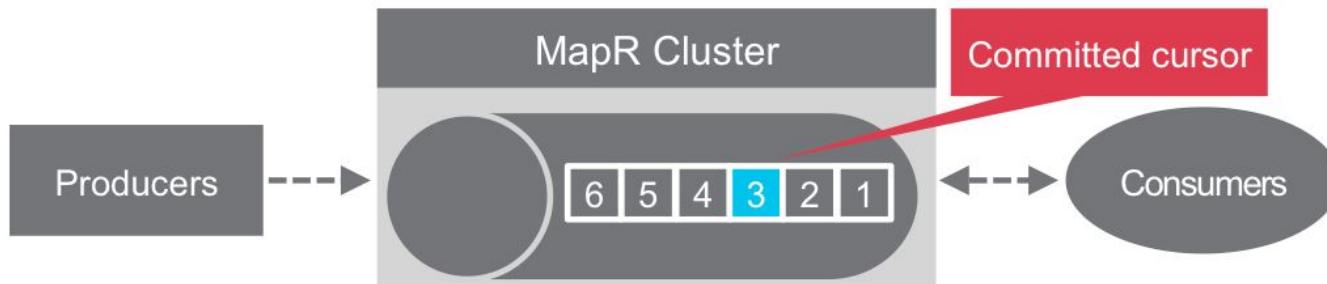


## Consumer Properties: Cursor

Property	Description
<code>enable.auto.commit</code>	If true, periodically commits the highest offsets of the messages fetched by the consumer (default: true)
<code>auto.commit.interval.ms</code>	Frequency in milliseconds that the offsets are committed (default: 1000ms)

# Commit Cursor: Automatic or Manual?

`enable.auto.commit` – if cursors should be committed automatically (default true)



## Example: Consumer Properties

```
Properties props = new Properties();
props.put("enable.auto.commit", "false");
consumer = new KafkaConsumer<String, String>(props);
. . .
// manually commit offsets of messages received
consumer.commitSync();
```

# Goals



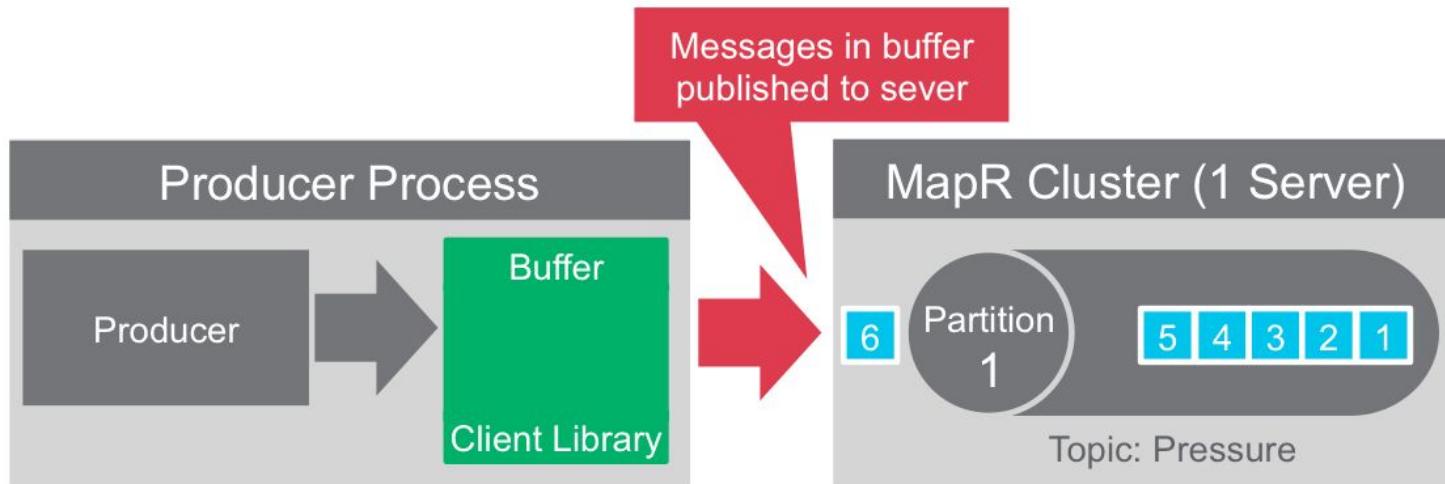
# Learning Goals

---

- Describe producer properties and options for:
  - buffering and batching of messages
  - publishing to partitions
- Describe consumer properties and options for:
  - fetching data
  - consumer groups
  - read cursors
- **Explain messaging semantics**

# Delivery Semantics: At Least Once

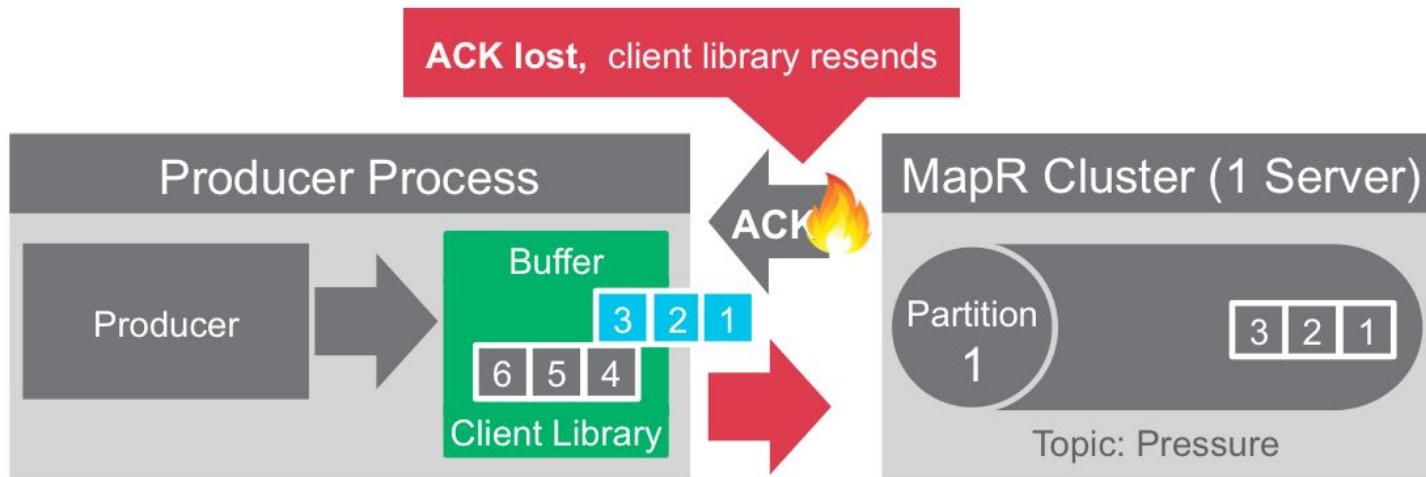
Once published to server, messages will be delivered



# Message Duplication: Producer Library Resends

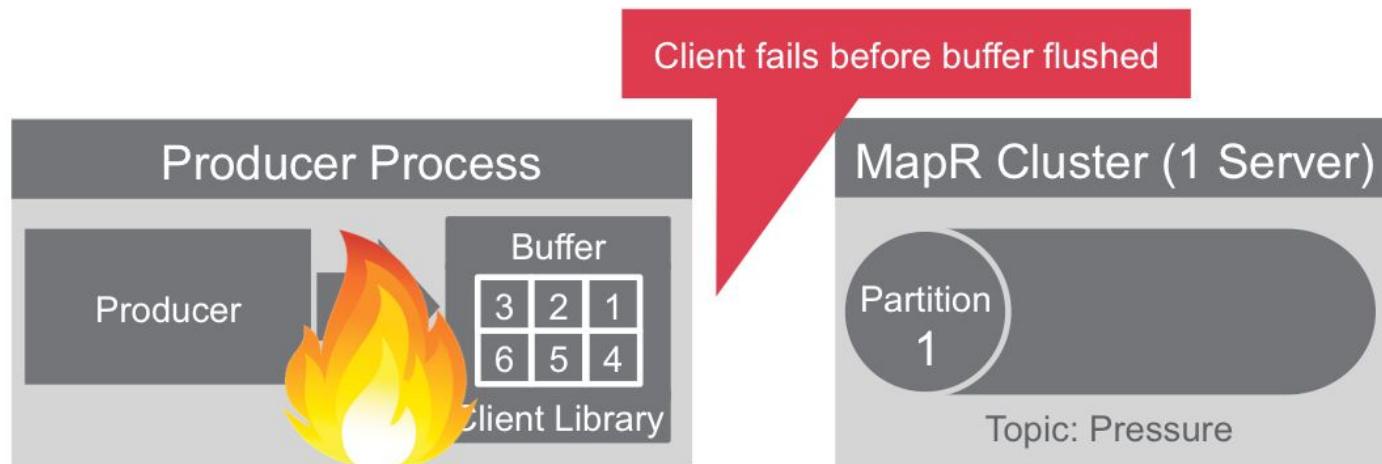
Messages may be duplicated

- producer side
- retransmit on network issue



# Avoid Losing Unpublished Messages

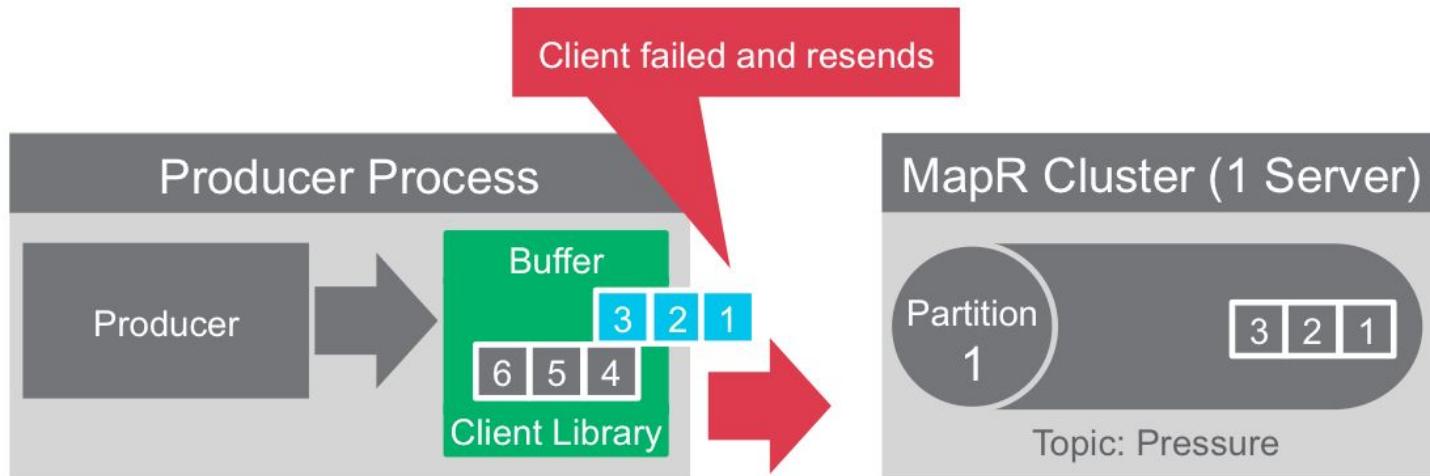
If a client fails before client buffer flushed, message unpublished producer can provide callback, resend on restart



# Message Duplication: Producer Resends

Messages may be duplicated

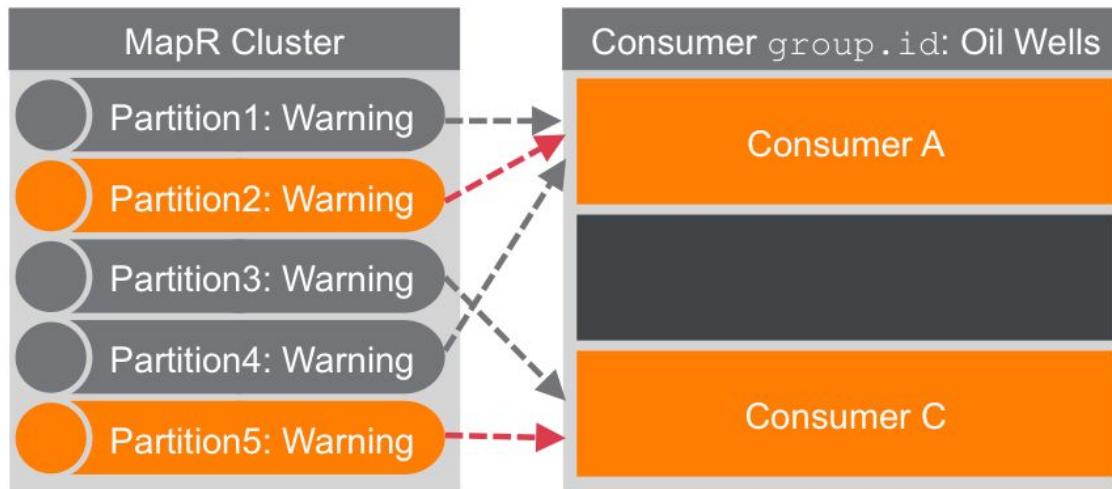
- Producer side
- Producer might fail and resend previously sent messages on restart



# Message Duplication: Consumer Cursor Not Committed

Messages may be duplicated

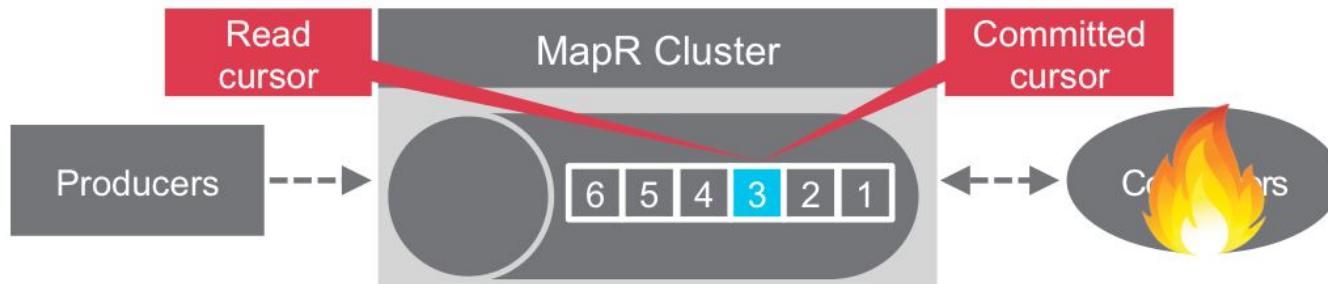
- consumer side
  - consumer partition re-assignment



# Messages Duplication: Consumer Fails Before Commit

Messages may be duplicated

- consumer side
  - client might crash after get but before `cursor.commit()`



# Delivery Semantics: Message Duplication

---

## Suggestions if duplicate messages are a problem:

- design your applications for idempotent messages
- add unique ID for duplicate detection